


Making Sense of Intelligence, both natural and artificial

Mark Ho

New York University
Department of Psychology


1




Computation and
Decision-Making
Lab

Making Sense of Intelligence



Psychology / Neuroscience



Purple Red Brown
Red Green Blue




Machine Learning / Robotics




Making Sense of Intelligence

Psychology / Neuroscience




Purple Red Brown
Red Green Blue



Engineering
Artificial Intelligence

Making Sense of Intelligence





Making Sense of Intelligence




Making Sense of Intelligence Hydraulics?

Natural Hydraulics





Artificial Hydraulics




Making Sense of Intelligence Hydraulics?

Natural Hydraulics



Artificial Hydraulics



Making Sense of Intelligence Hydraulics?

Natural Hydraulics

Fluid mechanics

[Article](#) [Talk](#)

From Wikipedia, the free encyclopedia

Fluid mechanics is the branch of [physics](#) concerned with the [mechanics of fluids](#) (liquids, gases, and plasmas) and the [forces](#) on them.^{[1]:3} It has applications in a wide range of disciplines, including [mechanical](#), [aerospace](#), [civil](#), [chemical](#) and [biomedical engineering](#), [geophysics](#), [oceanography](#), [meteorology](#), [astrophysics](#), and [biology](#).


$$\tau = -\mu \frac{du}{dn}$$

$$\tau_{ij} = \mu \left(\frac{\partial v_i}{\partial x_j} + \frac{\partial v_j}{\partial x_i} - \frac{2}{3} \delta_{ij} \nabla \cdot \mathbf{v} \right)$$

τ is the shear stress exerted by the fluid ("drag"),
 μ is the fluid viscosity—a constant of proportionality
 $\frac{du}{dn}$ is the velocity gradient perpendicular to the direction of flow

The Navier–Stokes equations describe the force balance at a point in a fluid. They are:^{[13][14][15][16]}

$$\frac{\partial \mathbf{u}}{\partial t} + (\mathbf{u} \cdot \nabla) \mathbf{u} = -\frac{1}{\rho} \nabla P + \nu \nabla^2 \mathbf{u} + \mathbf{f}$$





Two ways to make sense of intelligence: Rationality and computation

Rationales

- What problem is a system solving?
- *Functional* explanations
- Problem statements and utility functions



Computations

- How does a system find a solution?
What does it use to find a solution?
- *Mechanistic* explanations
- Algorithms and search strategies




Also see Tinbergen, 1963 for a similar perspective in ethology

Mechanistic explanations of vision





Mechanism of vision: light hits retina, causing neural firing, etc.

Functional explanations of vision



Functional explanations of vision






Kersten & Yuille, 2003

Function of vision: Identifying objects in environment

Functional accounts of higher-level and social cognition

- Natural and artificial intelligence have very different mechanisms, but they can share similar functions
 - Different mechanisms: neurons versus silicon
 - Same function: detecting objects
- Function(s) of perception are fairly clear-cut
 - Even so, computational theory is very useful!
- Function(s) of higher-level and social cognition are less obvious
 - Goals, intentionality, agency, learning, curiosity, adaptation, intelligence, habits, communication, pedagogy, norms, cooperation, morality ...
 - Computational theory is extremely useful, maybe unavoidable


Functional accounts of higher-level and social cognition




Human Problem Solving




Human Problem Solving



Tolman (1932)




Miller, Galanter & Pribram (1960)



Newell & Simon (1972)

Planning


Standard Model of Planning (e.g., Heuristic Search)




Newell & Simon, 1976; Puterman, 1994; Sutton & Barto, 1998; 2018

Planning

Full Representation




Expert Eye-Movements




Tichomirov & Poznyanskaya (1966)

Planning

Full Representation




Expert Eye-Movements



Tichomirov & Poznyanskaya (1966)

Planning




How do humans represent tasks?

(Newell & Simon, 1972; Ohlsson, 2012)

People construct **value-guided construals**


Maze Navigation Task

Mazes constructed out of fixed walls (black)
and changing obstacles (blue)




Ho, Abel, Correa, Littman, Cohen & Griffiths (2022). *Nature*.


Maze Navigation Task



People plan and then act





Mean Awareness Responses




Standard Model of Planning

(e.g., Heuristic Search)



Value-Guided Construal



Value-Guided Construal


Value of Construal

$$V(c) = U(\pi_c) - C(c)$$

↑
Construal

Behavioral Utility

Cognitive Cost



Value-Guided Construal

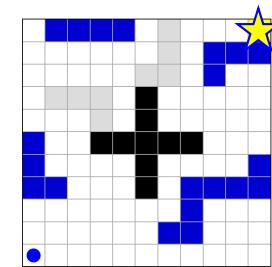

Value of Construal

$$V(c) = U(\pi_c) - C(c)$$

Behavioral Utility

Cognitive Cost

Construal



Value-Guided Construal

Value of Construal


$$V(c) = U(\pi_c) - C(c)$$

Behavioral Utility


Cognitive Cost

↑
Plan Computed with Construal


Construal



Computed Plan



X Utility of
Computed Plan




Value-Guided Construal

Value of Construal

Behavioral Utility


Cognitive Cost

$$V(c) = U(\pi_c) - C(c)$$




Value-Guided Construal

$$V(c) = U(\pi_c) - C(c)$$




Computed Plan



Complexity of construal

Utility of
Computed Plan



Value-Guided Construal


Value of Construal

Behavioral Utility

Cognitive Cost




$$V(c) = U(\pi_c) - C(c)$$

Construal




Value-Guided Construal

Value of Construal	Behavioral Utility	Cognitive Cost
	$V(c) = U(\pi_c) - C(c)$	





Do people form optimal construals?




*Pre-registered model predictions: <https://osf.io/zpq69>

Do people form optimal construals?



Obstacles split by 0.5 awareness: $\chi^2(1) = 23.03, p = 1.6 * 10^{-6}, w = 0.52$


Crossing a Bridge




Crossing a Bridge





Critical Mazes and Obstacles



Obstacle Recall Probe



Critical Mazes and Obstacles



Two concerns


1. Are patterns of construal a side-effect of other perceptual or cognitive mechanisms?

2. Is value-guided construal computationally feasible?

Is construal a side-effect?


Value-Guided Construal

Construal actively shapes planning




Planning **without** Construal

Construal is *side-effect* of other processes
(e.g., low-level perceptual cues)



Is construal a side-effect of other processes?

Value-guided Construal




Is construal a side-effect of other processes?

Value-guided Construal

Trajectory-based Heuristic Search

Barto, Bradtke & Singh, 1995
Bonet & Geffner, 2003

Real-Time Dynamic Programming
(RTDP) + Heuristic




Is construal a side-effect of other processes?

Value-guided Construal

Trajectory-based Heuristic Search

Barto, Bradtke & Singh, 1995
Bonet & Geffner, 2003

Real-Time Dynamic Programming
(RTDP) + Heuristic



Is construal a side-effect of other processes?

Value-guided Construal


Trajectory-based Heuristic Search

Barto, Bradtke & Singh, 1995
Bonet & Geffner, 2003

Graph-based Heuristic Search

Hart, Nilsson & Raphael, 1968
Hansen & Zilberstein, 2001

LAO* + Heuristic
(A* for MDPs)



Is construal a side-effect of other processes?

Value-guided Construal

Trajectory-based Heuristic Search

Barto, Bradtke & Singh, 1995
Bonet & Geffner, 2003

Graph-based Heuristic Search

Hart, Nilsson & Raphael, 1968
Hansen & Zilberstein, 2001



Related to ***experience replay***

Mattar & Daw, 2018; Pfeiffer & Foster, 2013;
Diba & Buzsáki, 2007


Is construal a side-effect of other processes?

Value-guided Construal

Trajectory-based Heuristic Search

Graph-based Heuristic Search

Is construal a side-effect?



Is construal a side-effect?


Value-guided Construal
Trajectory-based Heuristic Search
Graph-based Heuristic Search
Bottleneck Distance
Successor Representation Overlap
Minimal Path Distance
Minimal Path Distance Step
Goal Distance
Start Distance
Wall Distance
Center Distance

Step 1: Fit single global model to recall responses

Step 2: Remove each one, refit

Step 3: Compare

Addressing concerns 1 and 2: *Is construal a side-effect?*



Two concerns

1. Are patterns of construal a side-effect of other perceptual or cognitive mechanisms?

Even when accounting for alternative factors, value-guided construal explains responses

2. Is value-guided construal computationally feasible?

Is construal computationally feasible?


Value of Construal

Behavioral Utility



Cognitive Cost

$$V(c) = U(\pi_c) - C(c)$$


Can construals be efficiently optimized *in principle*?



One approach: Construal search



One approach: Construal search



Two concerns

1. Are patterns of construal a side-effect of other perceptual or cognitive mechanisms?



Even when accounting for alternative factors, value-guided construal explains responses

2. Is value-guided construal computationally feasible?

Optimal construal could be efficiently approximated using search-based methods.

Functional Fixedness


Goal: Mount the candle to the wall and light it



Participants used the wrong task representation
(e.g., construing the box as a container and not as a support)

Duncker, 1945


Functional Fixedness: A consequence of construal




Ho, Cohen & Griffiths (in press) *Psychological Science*.

Blocks and Notches

True maze




Blocks-only
construal of maze




Blocks and Notches

True maze




Blocks-only
construal of maze




Blocks and Notches


True maze




Blocks and notches
construal of maze




Notch Necessary Maze




Notch Unnecessary Maze




Test Maze




Blocks-and-Notches
Construal and Planning




Blocks-only
Construal and Planning




Blocks-and-Notches
Construal and Planning



Blocks-only
Construal and Planning






Training Condition: $\beta = -1.90$, SE = 0.17, $z = -11.22$, $p < 2 \times 10^{-16}$

Trial Number: $\beta = 0.11$, SE = 0.02, $z = 5.08$, $p = 3.7 \times 10^{-7}$

Interaction: $\beta = 0.09$, SE = 0.04, $z = 2.15$, $p = 0.03$


Modeling Functional Fixedness



Modeling Functional Fixedness

Switching

is **costly**




Complexity and switching best explain behavior (i.e. functional fixedness)


Model	df	AIC	ΔAIC
No complexity or switch cost	1	132094	2804
Only complexity cost	2	130535	1245
Only switch cost	2	130057	767
Both complexity and switch cost	3	129289	0

Modeling Functional Fixedness

Switching
is **costly**



Complexity costs predict reaction time




Making sense of social intelligence

Communicative
Interactions



Ho et al. (2021) *JEP:G*;
Ho et al. (2018) *Cog Sci*;
Ho et al. (2016) *NeurIPS*

Coordination and
Joint Planning



Ho et al. (2016) *CogSci*;
Carroll, et al. (2019) *NeurIPS*.

Natural stupidity and artificial intelligence

My colleagues, they study artificial intelligence; me, I study natural stupidity.

- Amos Tversky



How can we develop AI systems that complement,
rather than enhance, our natural stupidity?

Closing thoughts

- Making sense of intelligence
 - Rational vs mechanistic accounts of natural and artificial intelligence
- Task representations in human problem solving
 - Value-guided construal
 - Functional Fixedness
- Making sense of social intelligence
- Using cognitive science to inform interactive ML

Collaborators



Carlos Correa
NYU



David Abel
DeepMind



Tom Griffiths
Princeton




Michael Littman
Brown



Jon Cohen
Princeton



NYU



Computation and
Decision-Making
Lab

JOHN
TEMPLETON
FOUNDATION
Inspiring Awe & Wonder




TOYOTA
RESEARCH INSTITUTE

Thanks for listening!



NYU



Computation and
Decision-Making
Lab



Questions?