

Learning from data (experience) without
labels (rewards, optimal actions).



Self-Supervised Agents: Exploring and Learning with Minimal Feedback

Benjamin Eysenbach

Assistant Professor of Computer Science

May 1, 2025



PRINCETON
UNIVERSITY

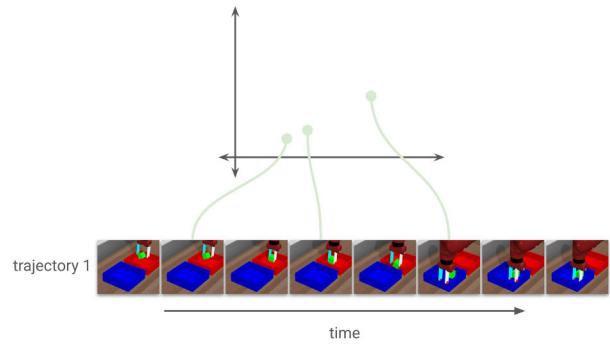
A thought experiment



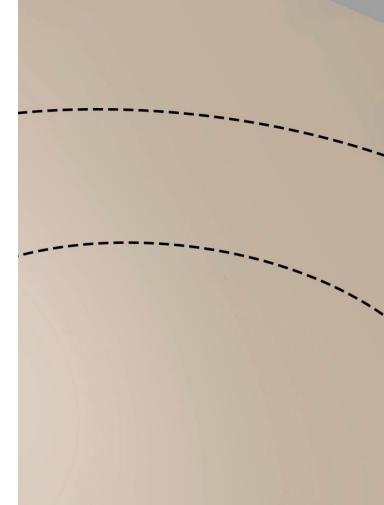
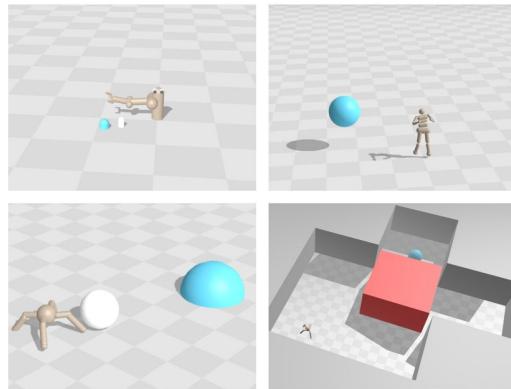
- Can a generative image model generate images different from those seen during training (e.g., a horse on the moon)?
- Can an LLM generate sentences different from those seen during training (e.g., write a poem about Tehran in the style of "twas the night before Christmas")?

👉 What is the right analogy for reinforcement learning (RL)?

Outline for today



1/ Self-supervised RL



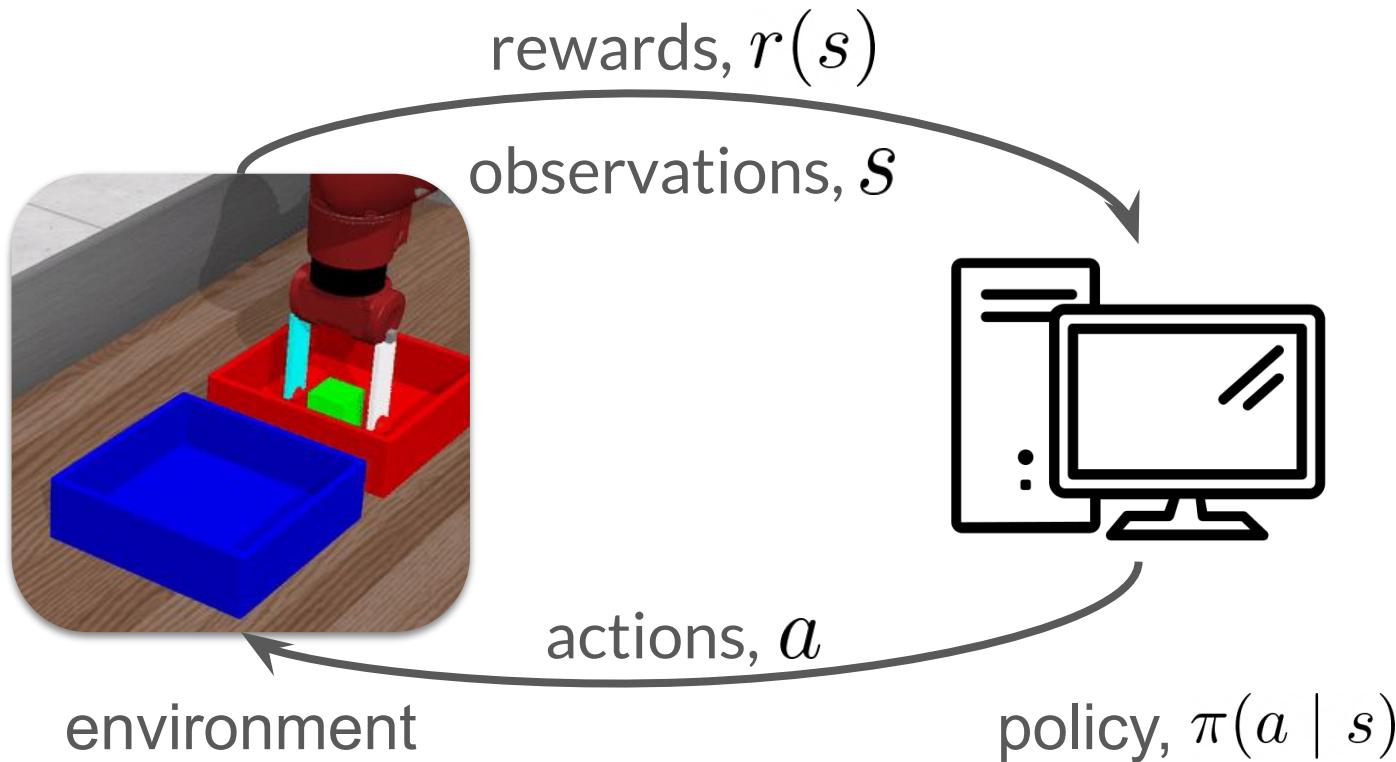
2/ Steps towards RL agents that can learn to do anything with minimal feedback

Papers highlighted:

1. BE, et al. *Contrastive learning as goal-conditioned reinforcement learning*. NeurIPS, 2022.
2. Bortkiewicz, et al. *Accelerating Goal-Conditioned RL Algorithms and Research*. ICLR, 2025.
3. Liu, Tang, BE. *A Single Goal is All You Need: Skills and Exploration Emerge from Contrastive RL without Rewards, Demonstrations, or Subgoals*. ICLR, 2025.
4. Myers, Ji, BE. *Horizon Generalization in Reinforcement Learning*. ICLR, 2025.

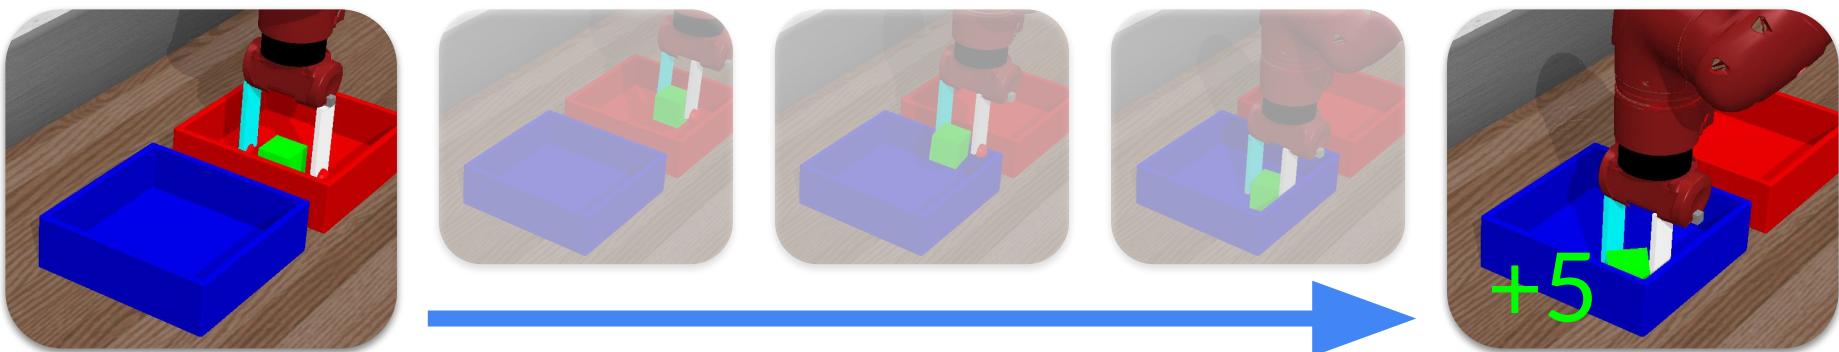
Background: What is RL?

Defining the reinforcement learning problem.



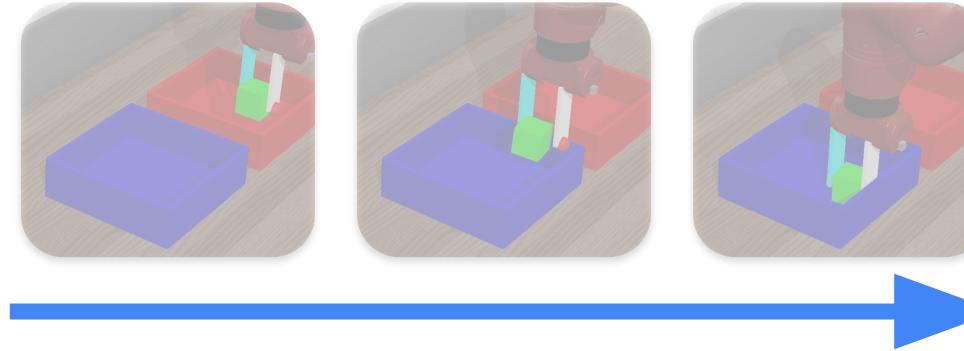
RL is hard because of limited feedback (rewards).

- 1 Only get feedback many steps into the future.



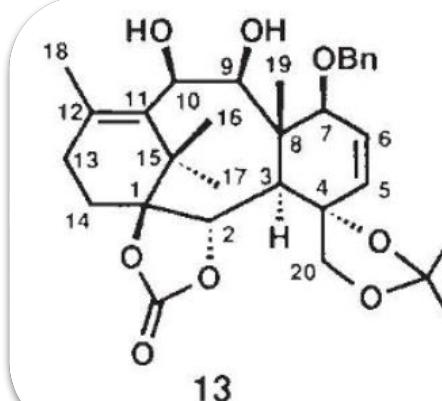
RL is hard because of limited feedback (rewards).

- 1 Only get feedback many steps into the future.
- 2 Limited feedback makes it challenging to learn from high-dimensional data.



RL is hard because of limited feedback (rewards).

- 1 Only get feedback many steps into the future.
- 2 Limited feedback makes it challenging to learn from high-dimensional data.
- 3 Practically, rewards and optimal actions are hard to s



```
def compute_reward(observation):
    objPos = obs[3:6]
    (rightFinger, leftFinger) = (self._get_site_pos('rightEndEffector'),
                                 self._get_site_pos('leftEndEffector'))
    fingerCOM = (rightFinger + leftFinger) / 2
    heightTarget = self._targetHeight
    placingGoal = self._targetPos
    reachDist = np.linalg.norm(objPos - fingerCOM)
    placingDist = np.linalg.norm(objPos[:2] - placingGoal[:2])
    def reachReward():
        reachRew = -reachDist
        reachDistxy = np.linalg.norm(objPos[:2] - fingerCOM[:2])
        zRew = np.linalg.norm(fingerCOM[2] - self.init_fingerCOM[2])
        if reachDistxy < 0.05:
            reachRew = -reachDist
        else:
            reachRew = -reachDistxy - zRew
        if reachRew > 0:
            reachRew = 1
        return reachRew
    def pickCompletionCriteria():
        tolerance = 0.01
        if objPos[2] >= heightTarget - tolerance:
            return True
        else:
            return False
    if pickCompleted():
        self.pickCompleted = True
    def objDropped():
        return objPos[2] < self.objHeight +
               > 0.02 and reachDist > 0.02
    def placeCompletionCriteria():
        if abs(objPos[0] - placingGoal[0]) <
           - placingGoal[1]) < 0.05 and
           abs(objPos[1] - placingGoal[1]) < 0.05 and
           abs(objPos[2] - placingGoal[2]) < 0.05:
            return True
        else:
            return False
    // Yu et al., "Meta-world: A benchmark and evaluation for reinforcement learning," CoRL, 2020
```



How can we get the benefits of RL without
the challenge of reward engineering?

Background: Self-Supervised Learning

Self-supervised learning in other areas of ML: Learning from high-dimensional data with limited feedback.

Computer Vision



NLP

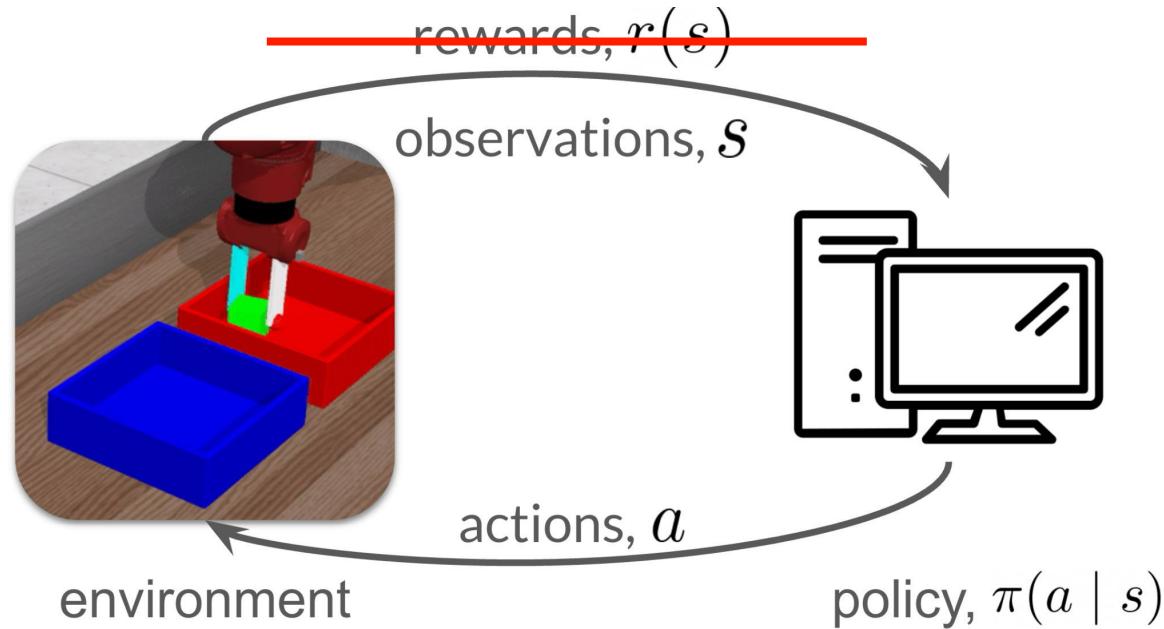
Give a step-by-step analysis of the worst-case time complexity of the bubble sort algorithm, with Python code examples, but write every sentence in the speaking style of a fast-talkin' wise guy from a 1940's gangster movie.



Alright, pal, listen up. The bubble sort algorithm, she's a real simple one. But the worst-case time complexity, that's a whole different story. You see, in the worst case, every element in the list gotta be swapped, so the time complexity is gonna be proportional to the square of the length of the list. Get it?

Self-Supervised Reinforcement Learning

Self-supervised RL



Goal-Conditioned RL: one type of self-supervised RL

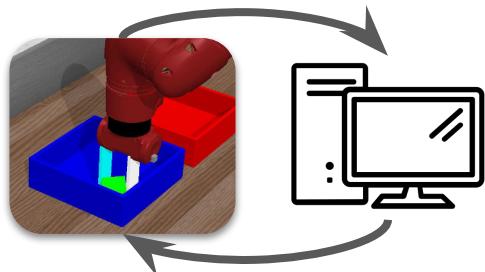
Inputs:

a) Online setting

OR

b) Offline setting

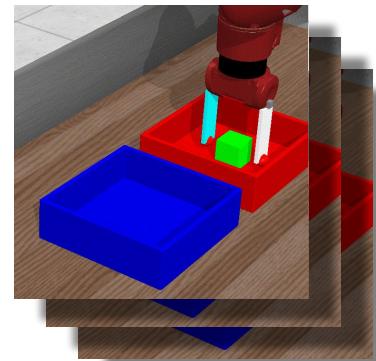
data: videos + actions



Output:

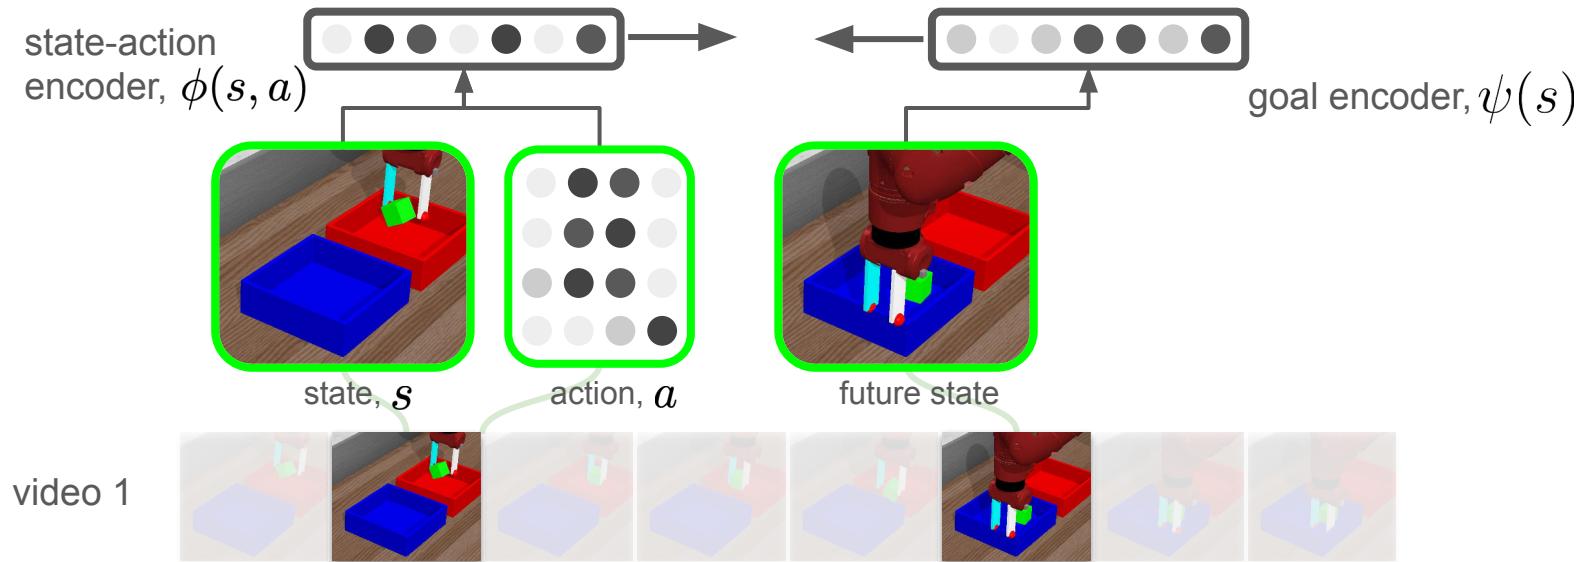
Goal-conditioned policy

$$\pi(a \mid s, g)$$



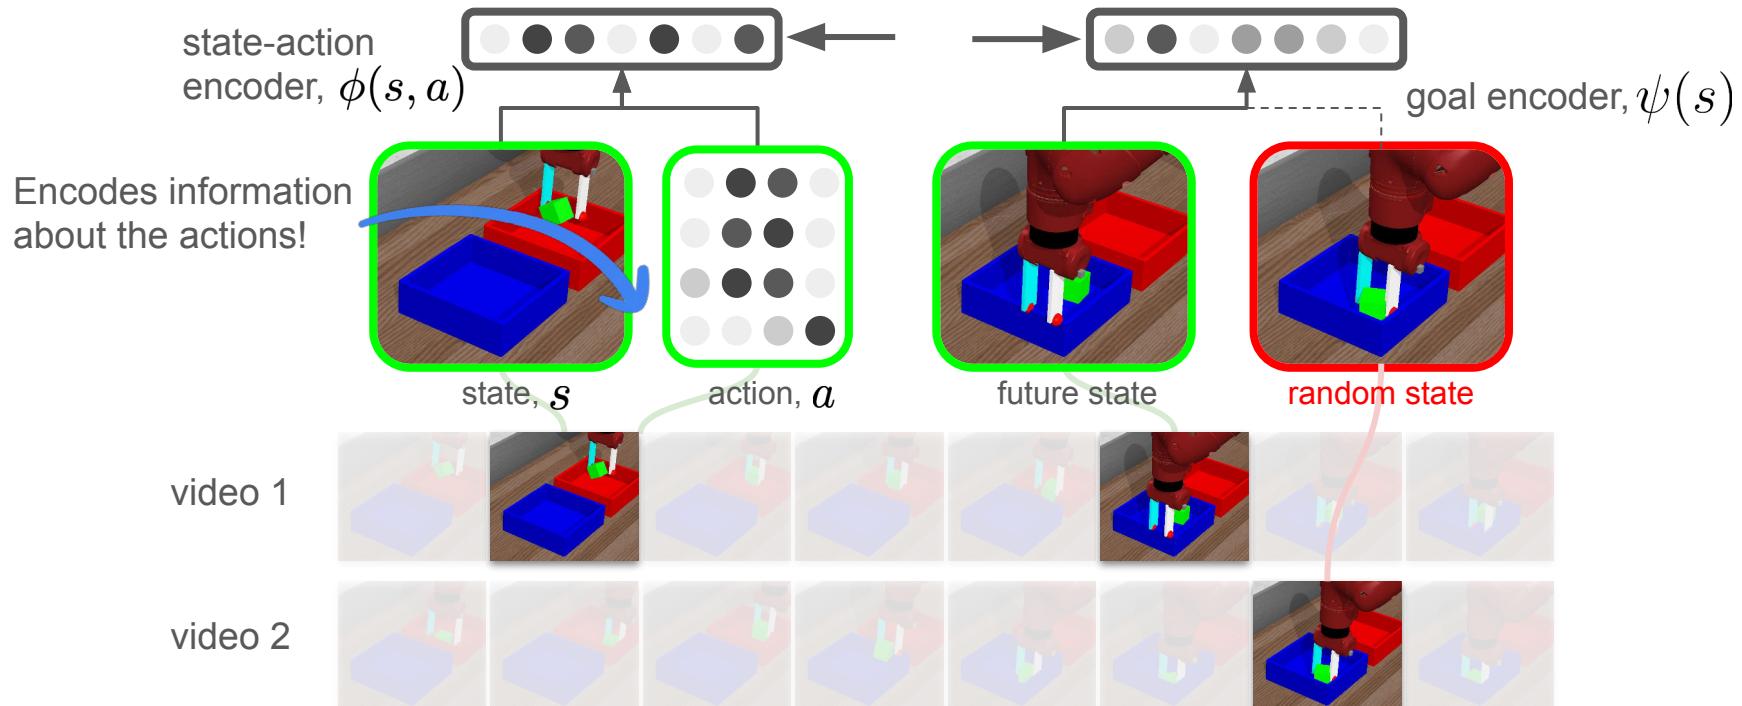
No reward or action labels required!

Learn representations via temporal contrastive learning



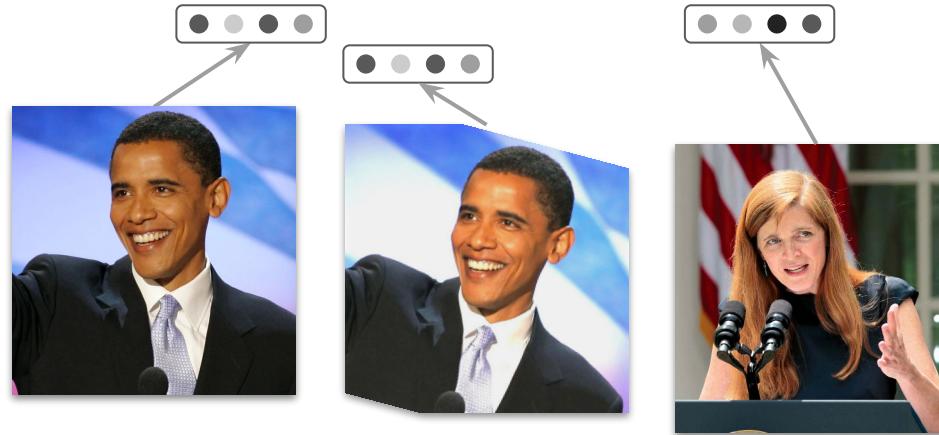
Much prior work in other domains: E.g., [Sermanet et al, 2018; van den Oord et al, 2018]

Learn representations via temporal contrastive learning



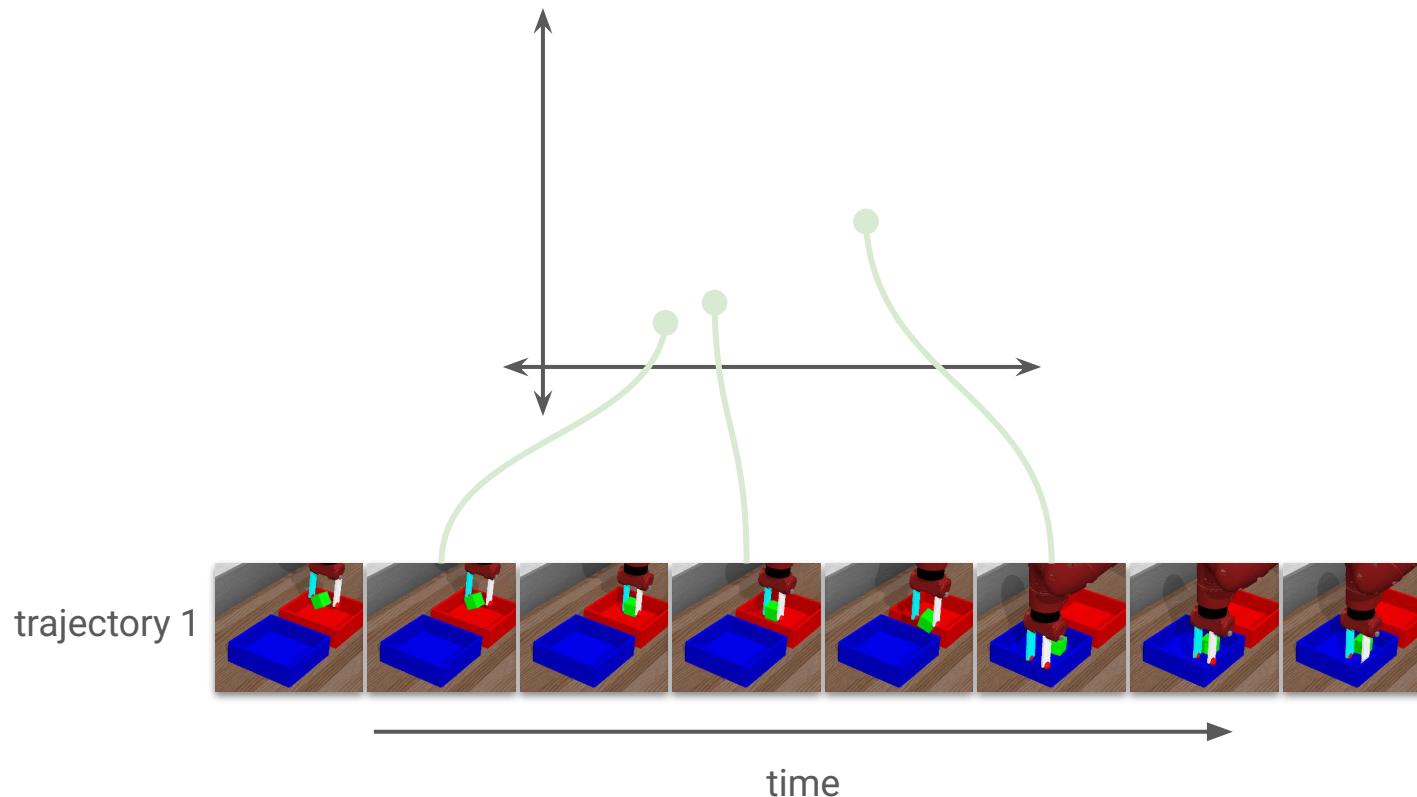
Much prior work in other domains: E.g., [Sermanet et al, 2018; van den Oord et al, 2018]

Compare with contrastive learning for vision



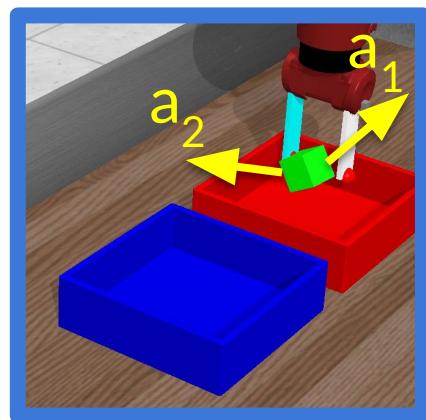
Augmentation: $p(x' | x)$

Intuition: representations encode temporal information

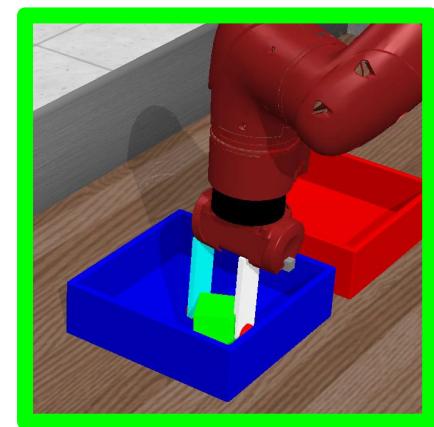
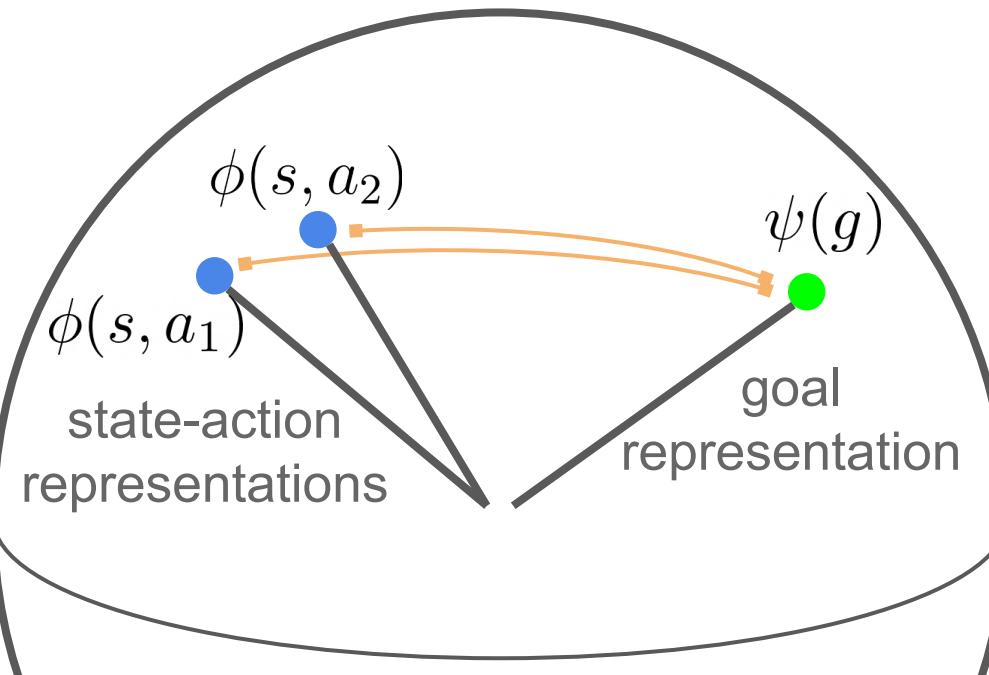


Use representations to extract goal-reaching skills.

$$\left\{ \pi(a | s, g) = \arg \max_a \phi(s, a)^T \psi(g), \dots \right\}$$

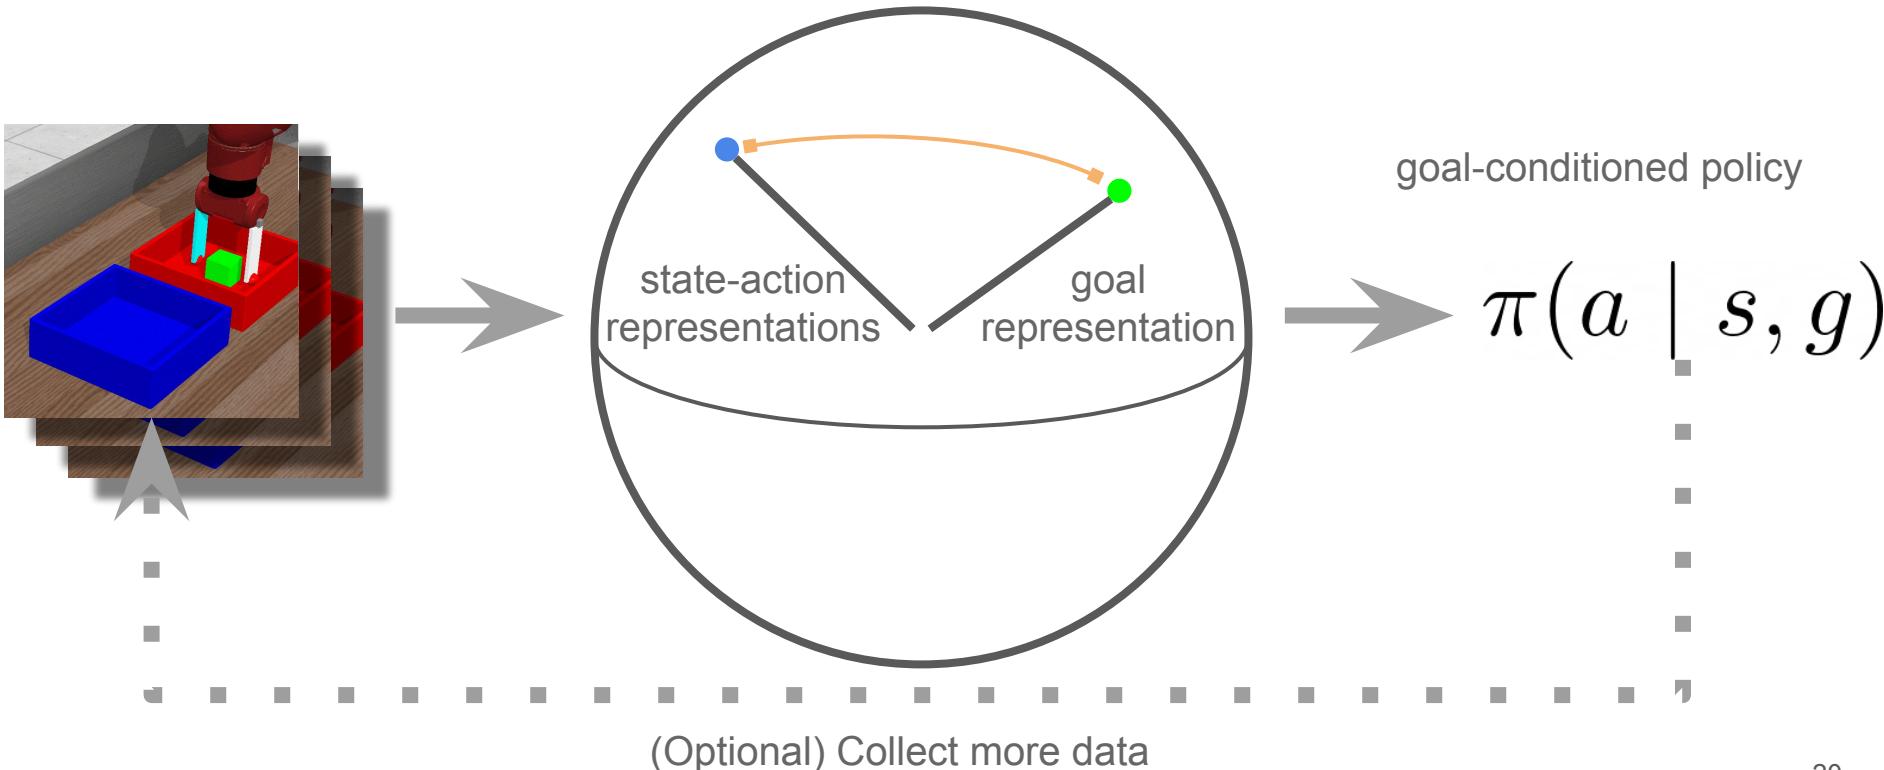


current state



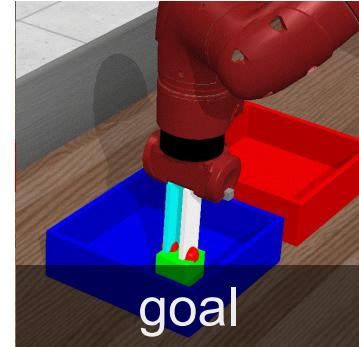
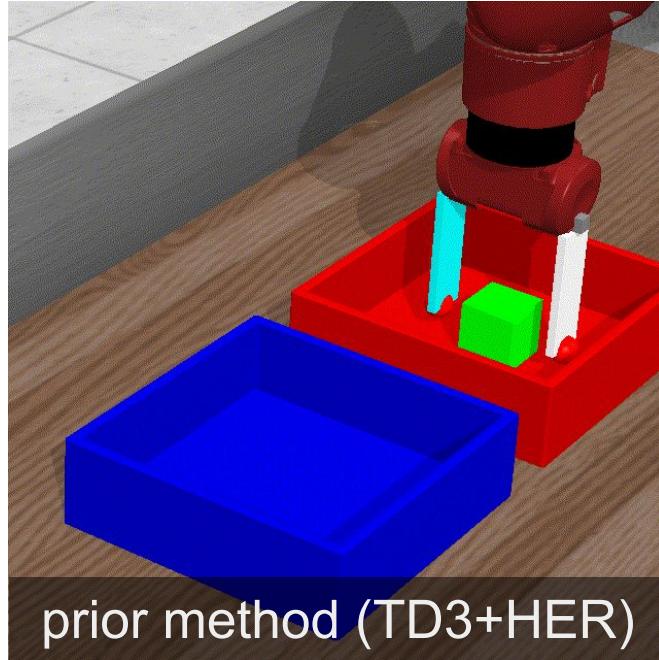
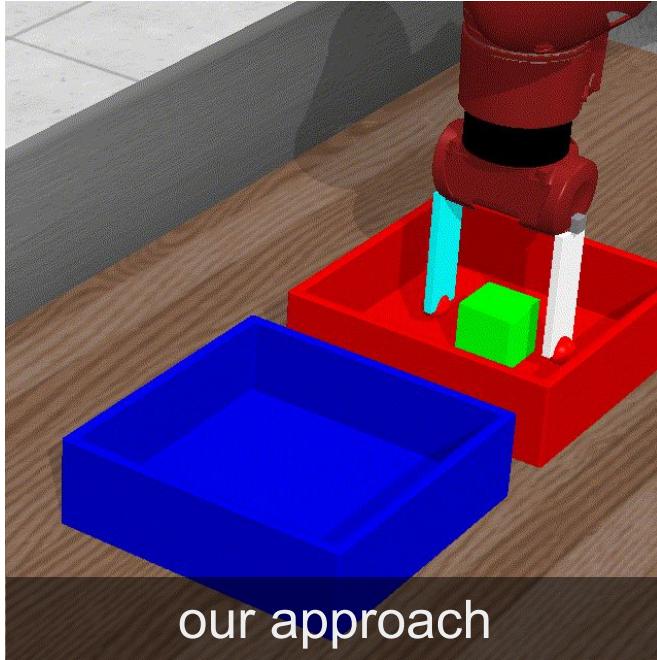
goal state

A complete method for learning (goal-reaching) skills from data.



Our complete method for learning (goal-reaching) skills from data.





self-supervised RL

RL

goal-conditioned RL

[Newell et al, 1961; Laird et al, 1983; Kaelbling, 1993; Andrychowicz et al, 2017.]

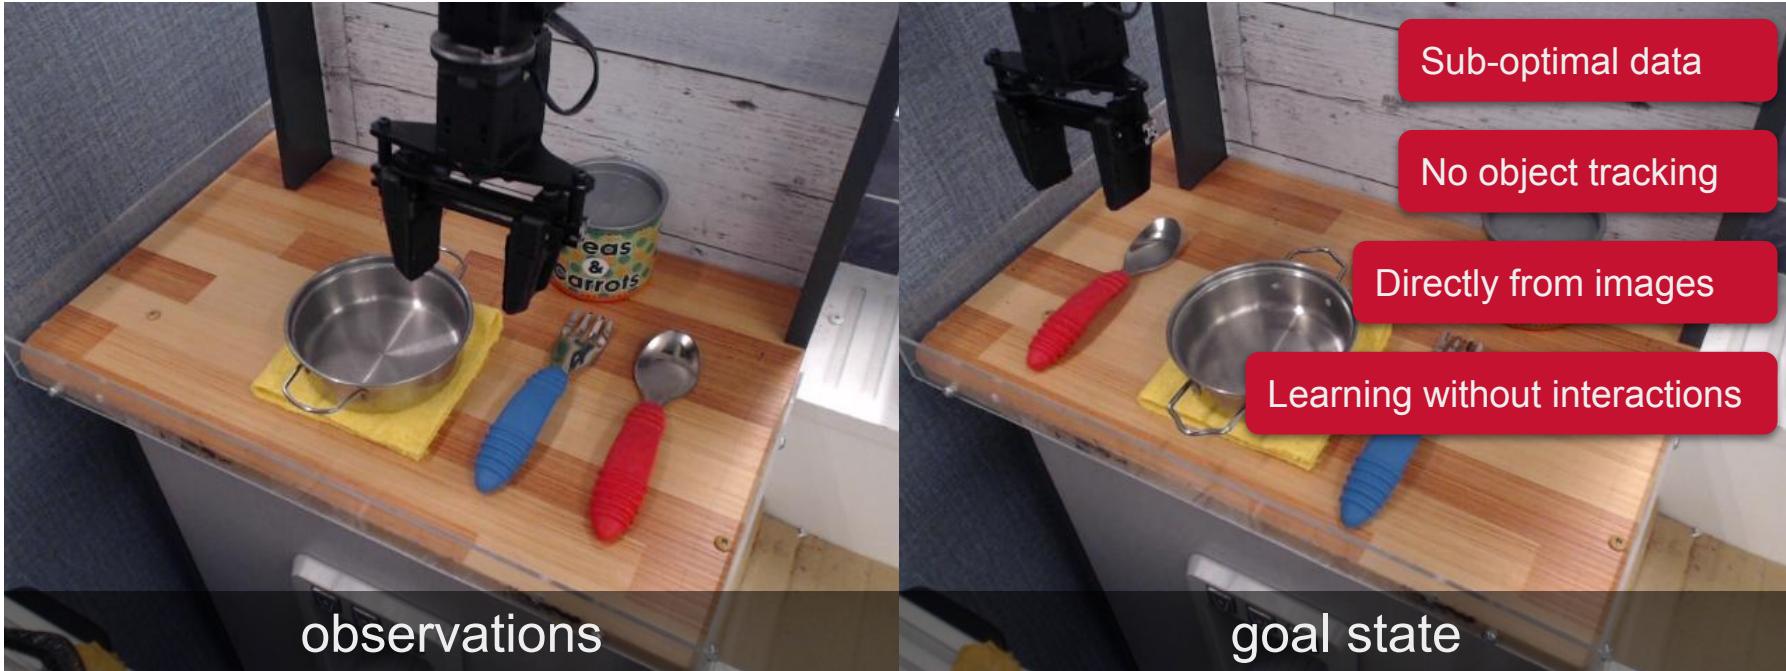
0

amount of rewards required

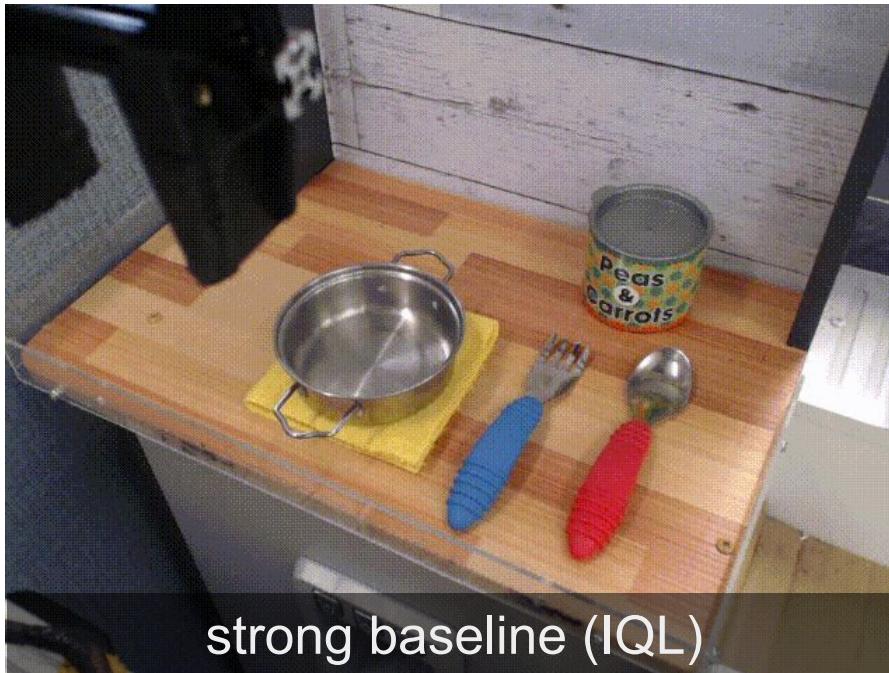
A couple real-world applications of
self-supervised RL

Evaluating the goal-reaching on a real robot.

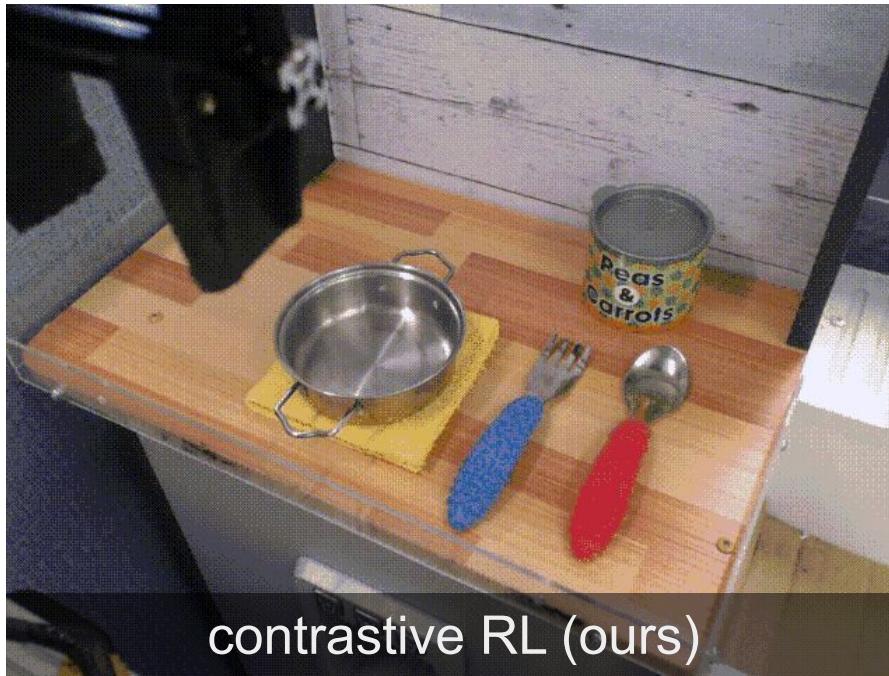
Chongyi Zheng,
Homer Walke



Evaluating the (goal-reaching) skills on a real robot.



Our goal-conditioned skills solve a real-world robotic task.



contrastive RL (ours)

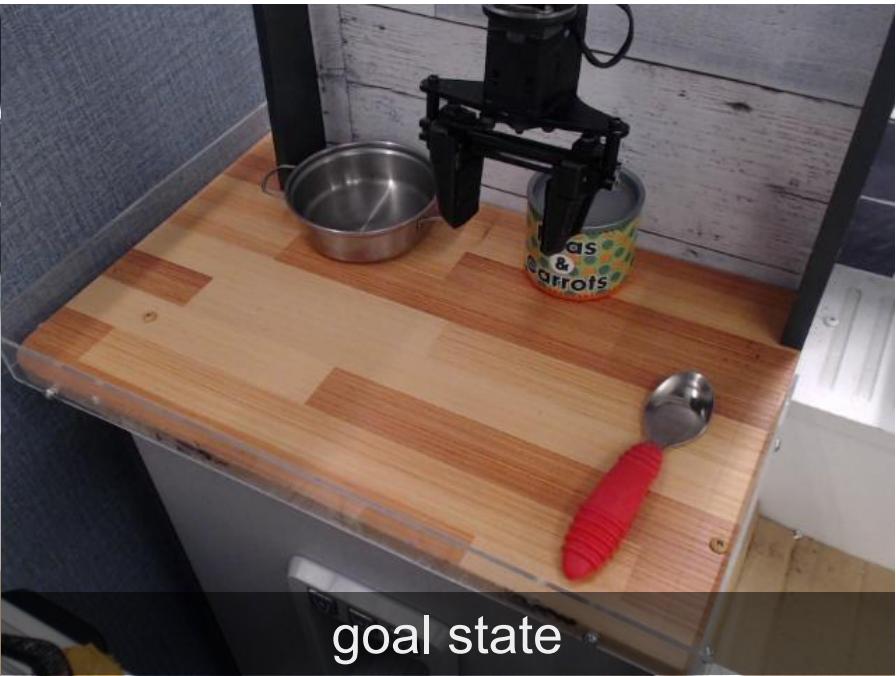


goal state

Solve new tasks by just "prompting" with a new goal

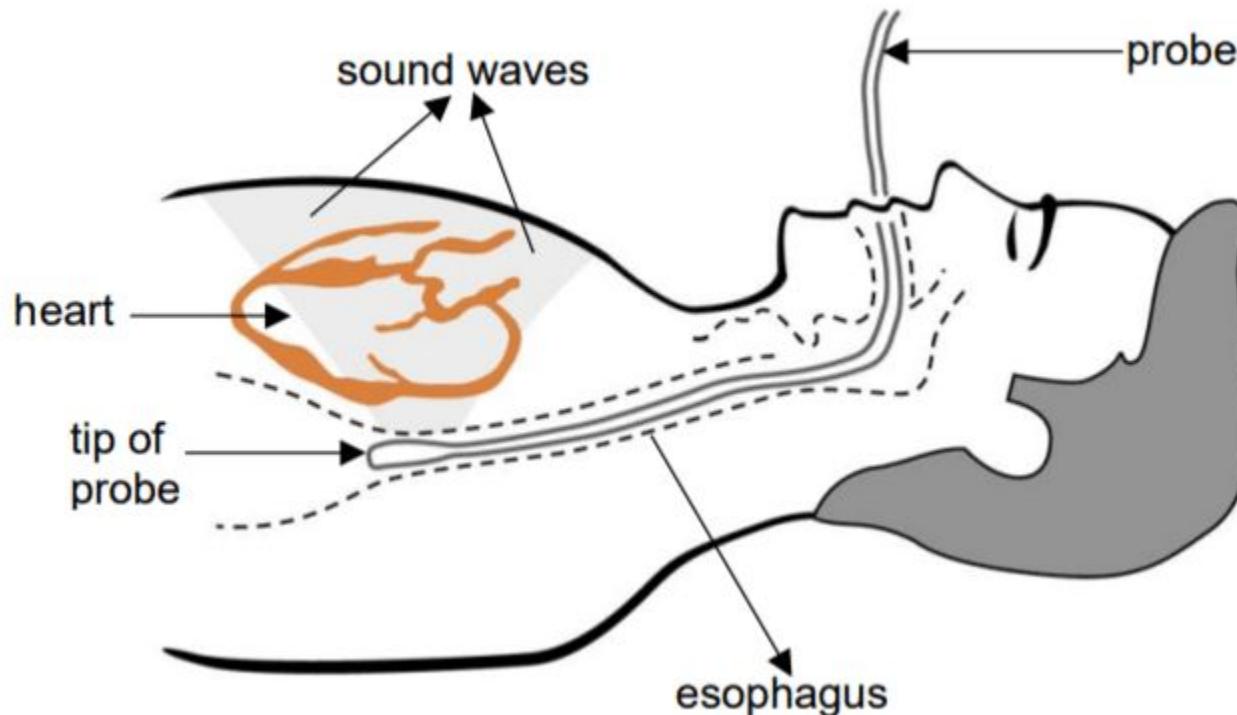


observations



goal state

Application to Esophageal Ultrasound

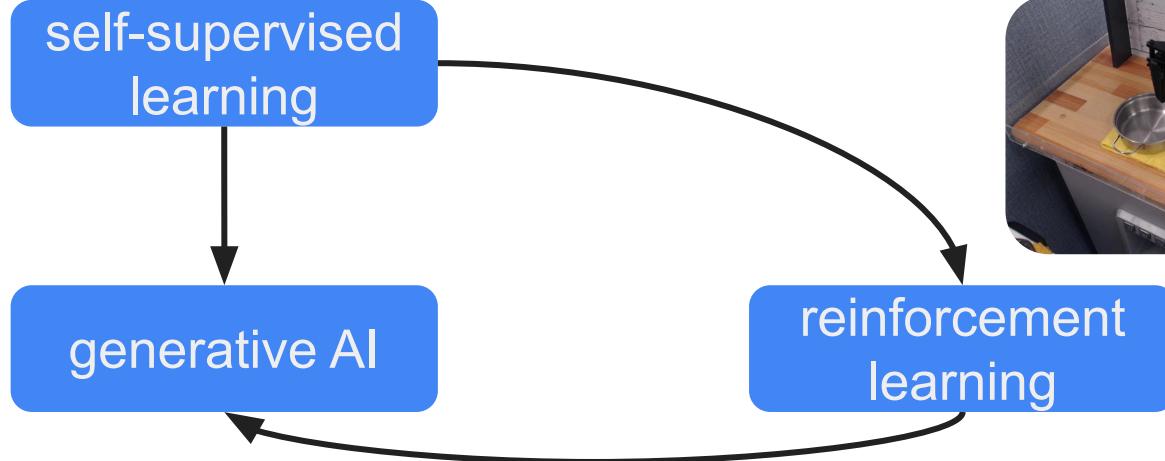


Amadou, Abdoul Aziz, et al. "Goal-conditioned reinforcement learning for ultrasound navigation guidance." *International Conference on Medical Image Computing and Computer-Assisted Intervention*. 2024.

How can we get the benefits of RL without the challenge of reward engineering?

```
def compute_reward(observation):
    objPos = observation[3:6]
    (rightFinger, leftFinger) = (self._get_site_pos('rightEndEffector'),
                                 self._get_site_pos('leftEndEffector'))
    fingerCOM = (rightFinger + leftFinger) / 2
    heightTarget = self.heightTarget
    placingGoal = self._target_pos
    reachDist = np.linalg.norm(objPos - fingerCOM)
    placingDist = np.linalg.norm(objPos[:2] - placingGoal[:2])
    def reachReward():
        reachRew = -reachDist
        reachDistxy = np.linalg.norm(objPos[:-1] - fingerCOM[:-1])
        zRew = np.linalg.norm(fingerCOM[-1] - self.init_fingerCOM[-1])
        if reachDistxy < 0.06:
            reachRew = -reachDist
        else:
            reachRew = -reachDistxy - zRew
        if reachDist < 0.05:
            reachRew = -reachDist + max(actions[-1], 0) / 50
        return (reachRew, reachDist)
    def pickCompletionCriteria():
        tolerance = 0.01
        if objPos[2] >= heightTarget - tolerance:
            return True
        else:
            return False
    if pickCompletionCriteria():
        self.pickCompleted = True
    def objDropped():
        return objPos[2] < self.objHeight + 0.005 and placingDist \
               > 0.02 and reachDist > 0.02
    def placeCompletionCriteria():
        if abs(objPos[0] - placingGoal[0]) < 0.05 and abs(objPos[1] \
               - placingGoal[1]) < 0.05 and objPos[2] < self.objHe
    else:
        return False
// Yu et al. "Meta-world: A benchmark and evaluation for multi-task reinforcement learning." CoRL, 2020
```





Self-supervised RL is generative AI

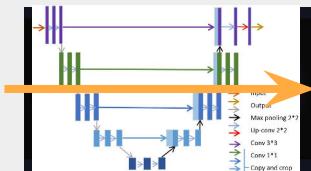
RL is generative modeling

Examples

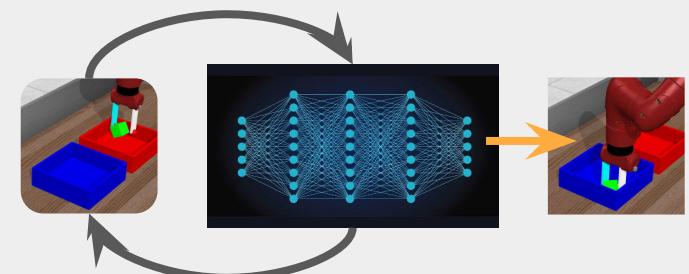
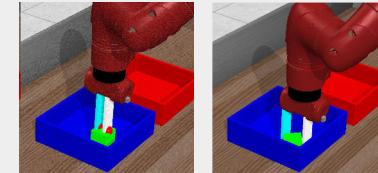
(Standard) Generative models



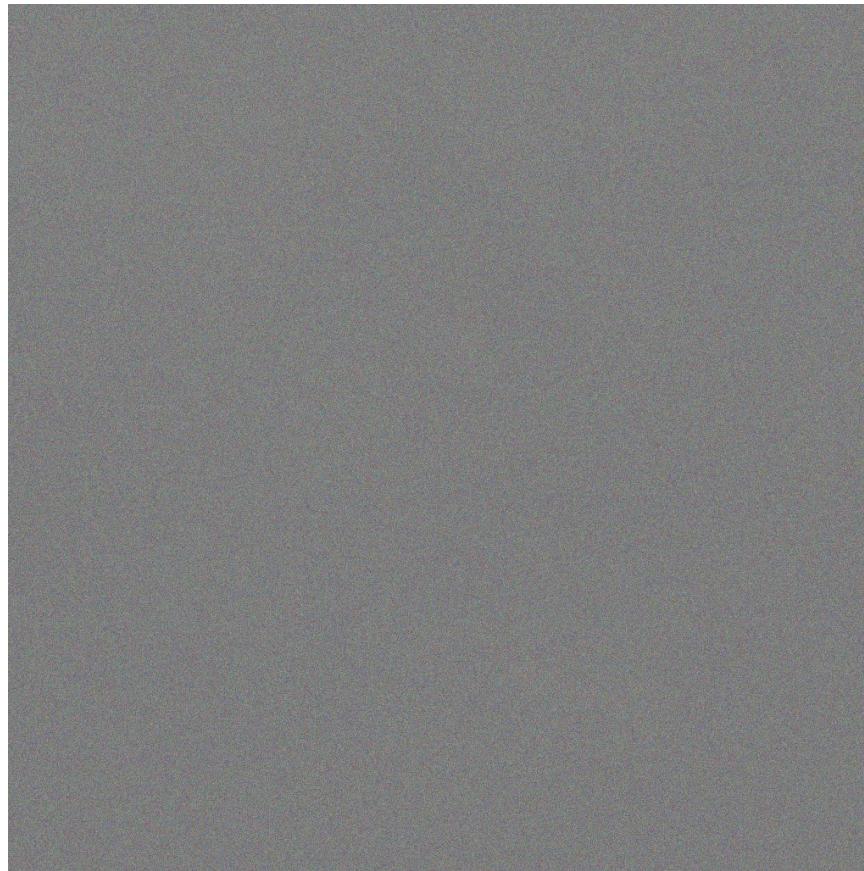
Sampling



Reinforcement Learning



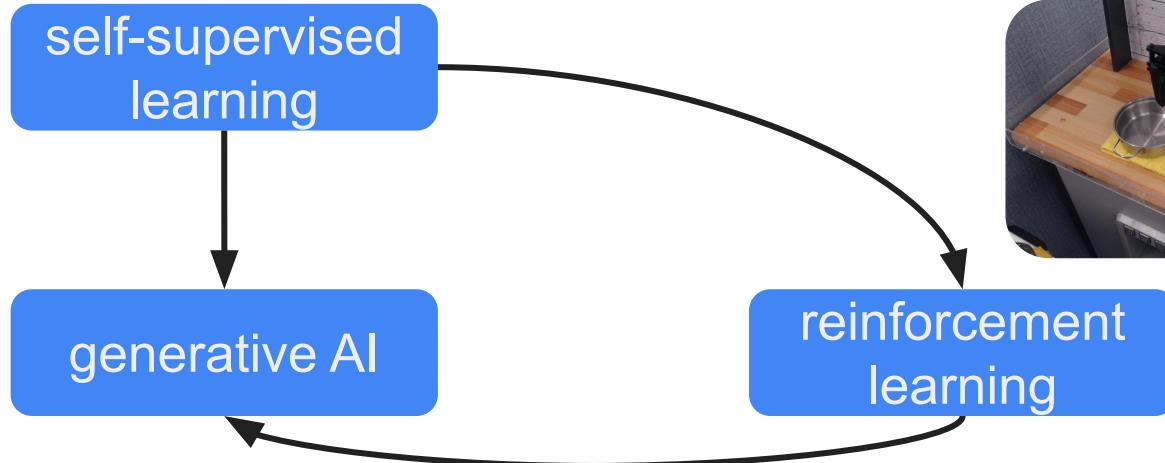
RL is generative AI: "a castle surrounded by mountains"



A robotic arm is shown in a 3D environment, placing colored blocks (pink, green, blue, yellow) to build a small castle structure. The arm is black and yellow.

This is your generative model.

"Build a castle surrounded by mountains"



Self-supervised RL is generative AI

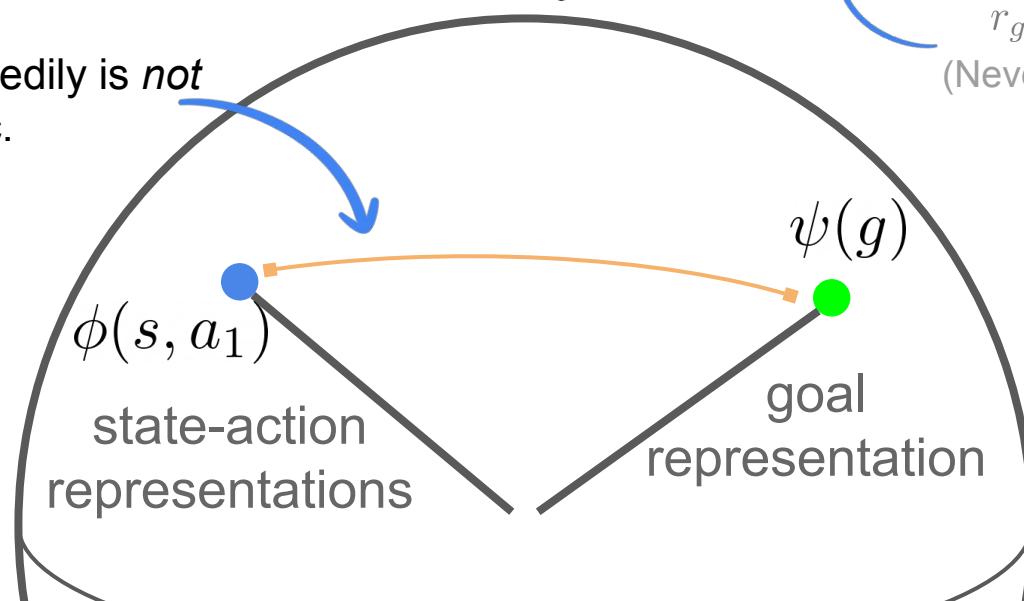
But what about tasks beyond goal reaching?

Theorem: The dot-product between the learned representations encodes the future returns, up to a constant.

$$e^{\phi(s,a)^T \psi(g)} = \frac{1 - \gamma}{p(g)} \mathbb{E}_\pi \left[\sum_t \gamma^t r_g(s_t, a_t) \right] = Q_g^\beta(s_t, a_t)$$

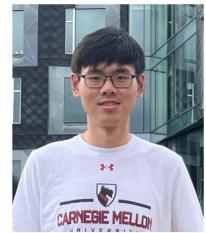
Acting greedily is *not* a heuristic.

$r_g(s, a) = p(s' = g \mid s, a)$
(Never need to compute this!)

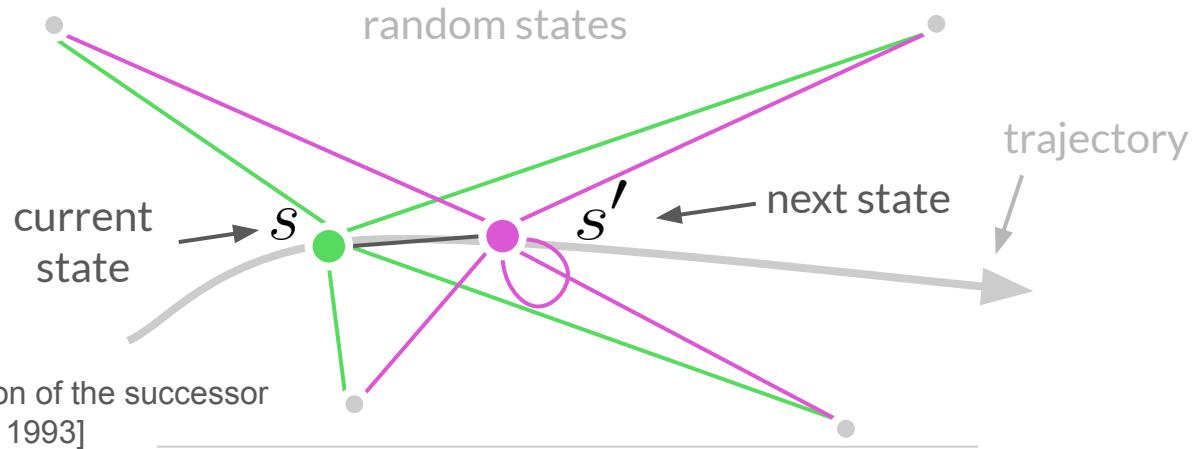


Estimating contrastive successor representations from off-policy data

Zheng, et al. "Contrastive Difference Predictive Coding." ICLR 2024.



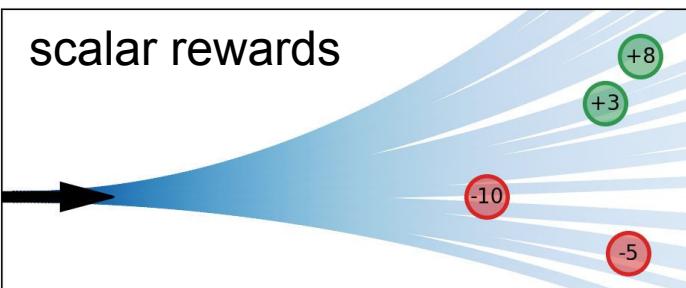
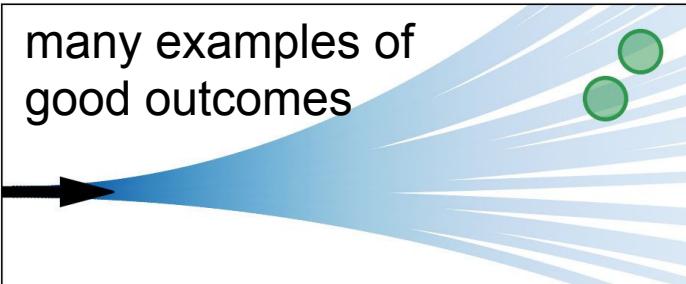
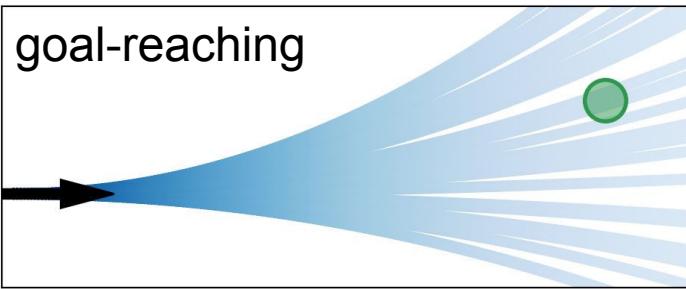
Chongyi Zheng



Contrastive approaches to different problems



Kyle Hatch, Bogdan Mazoure



Once you can estimate probability (ratios), you can solve many different types of problems

Mazoure, Bogdan, et al. *Contrastive value learning: Implicit models for simple offline RL*. CoRL, 2023.
Hatch, Kyle Beltran, et al. *Contrastive Example-Based Control*. L4DC, 2023

Steps towards RL agents that can do anything.

1. Fast simulators
2. Generalization

Lessons from generative AI in other domains

```
def compute_reward(observation):
    objPos = obs[3:6]
    (rightFinger, leftFinger) = (self._get_site_pos('rightEndEffector'),
                                 self._get_site_pos('leftEndEffector'))
    fingerCOM = (rightFinger + leftFinger) / 2
    heightTarget = self.heightTarget
    placingGoal = self.target_pos
    reachDist = np.linalg.norm(objPos - fingerCOM)
    placingDist = np.linalg.norm(objPos[:2] - placingGoal[:1])
    def reachReward():
        reachRew = -reachDist
        reachDistxy = np.linalg.norm(objPos[:-1] - fingerCOM[:-1])
        zRew = np.linalg.norm(fingerCOM[-1] - self.init_fingerCOM[-1])
        if reachDistxy < 0.06:
            reachRew = -reachDist
        else:
            reachRew = -reachDistxy - zRew
        if reachDist < 0.05:
            reachRew = -reachDist + max(actions[-1], 0) / 50
        return (reachRew, reachDist)
    def pickCompletionCriteria():
        tolerance = 0.01
        if objPos[2] >= heightTarget - tolerance:
            return True
        else:
            return False
    if pickCompletionCriteria():
        self.pickCompleted = True
    def objDropped():
        return objPos[2] < self.objHeight + 0.05 and
               > 0.02 and reachDist > 0.02
    def placeCompletionCriteria():
        if abs(objPos[0] - placingGoal[0]) < 0.05 and
           abs(objPos[1] - placingGoal[1]) < 0.05 and
           reachRew > 0.05:
            return True
        else:
            return False
    // Yu et al. "Meta-world: A benchmark and evaluation of reinforcement learning." CoRL, 2020
```



Self-supervised feedback



Lots of data



Lots of compute

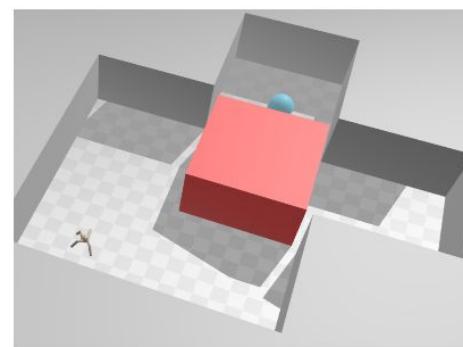
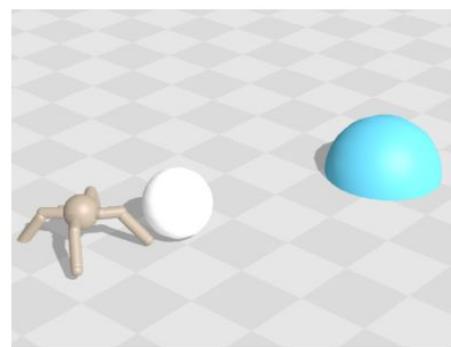
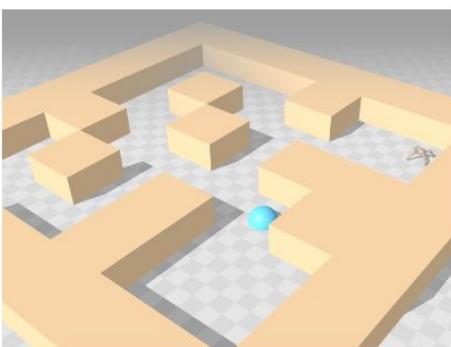
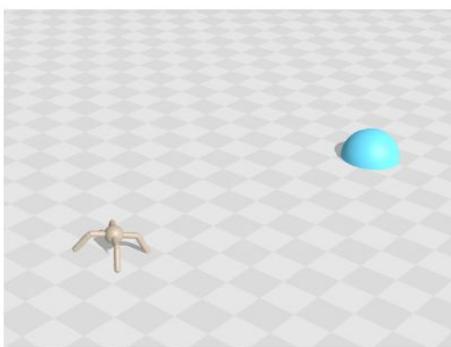
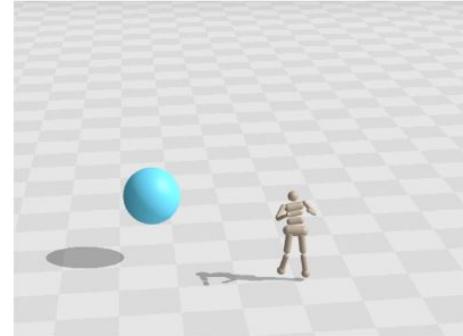
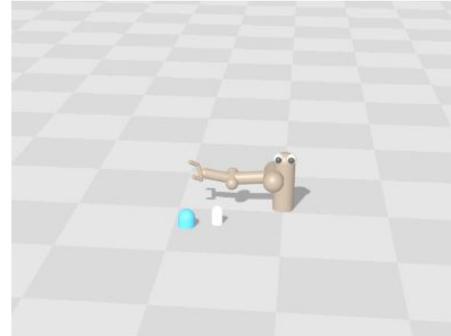
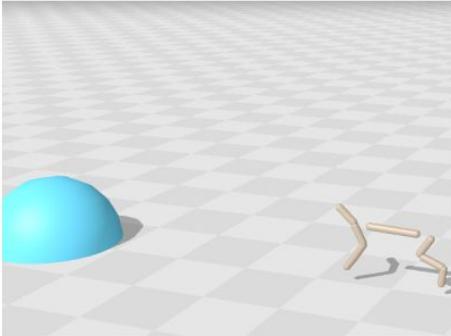
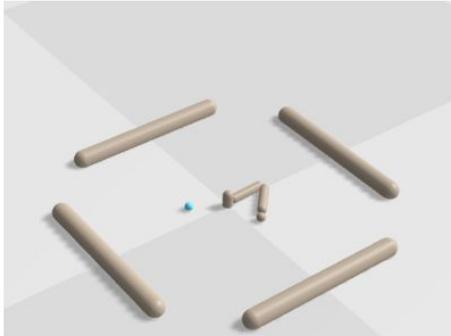


Standardized
benchmarks

Michał Bortkiewicz,
Władysław Pałucki



Fast Simulators: Jax GCRL



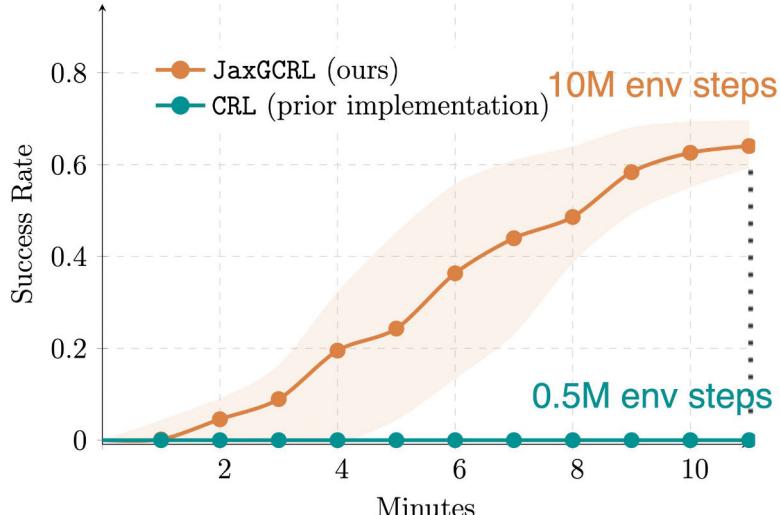
Bortkiewicz, Michał, et al. "Accelerating Goal-Conditioned RL Algorithms and Research."
ICLR, 2025

<https://github.com/MichalBortkiewicz/JaxGCRL/>



Fast Simulators: Jax GCRL

Fast Training

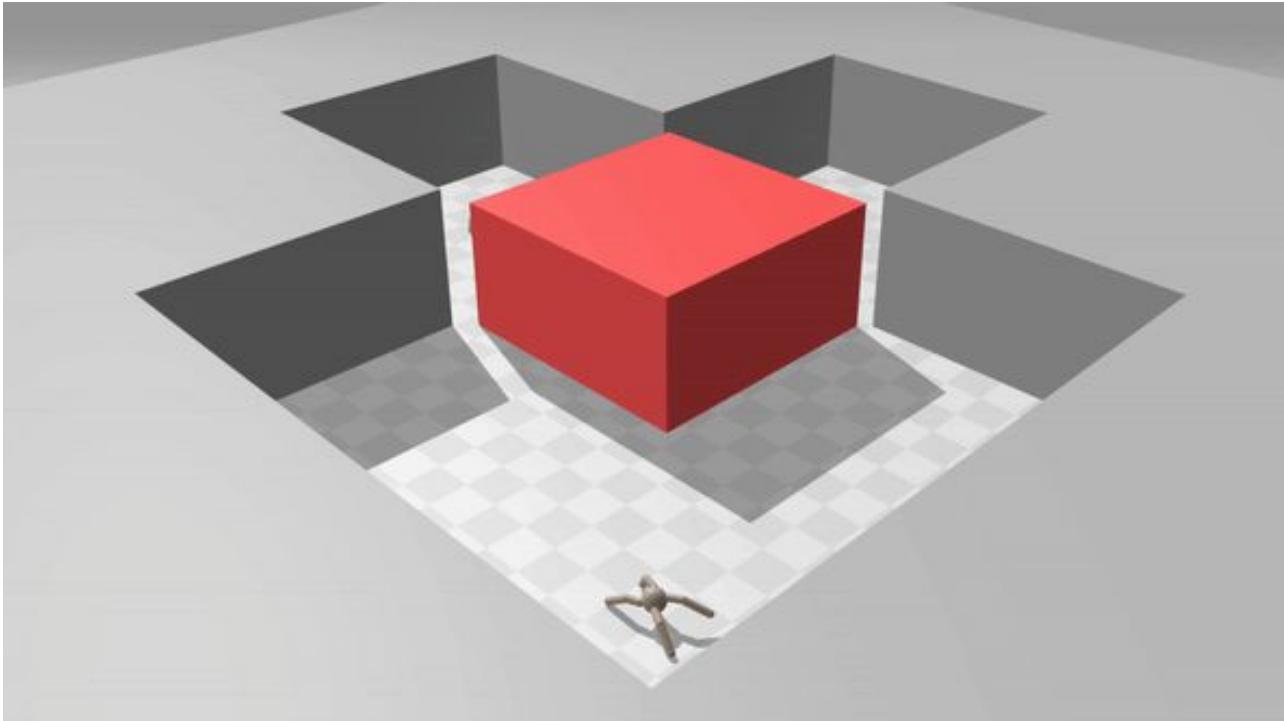


>20 times faster training!

Michał Bortkiewicz,
Władysław Pałucki



Fast Simulators: Jax GCRL



Bortkiewicz, Michał, et al. "Accelerating Goal-Conditioned RL Algorithms and Research." ICLR, 2025.

<https://github.com/MichalBortkiewicz/JaxGCRL/>

Michał Bortkiewicz,
Władysław Pałucki



Fast Simulators: Jax GCRL



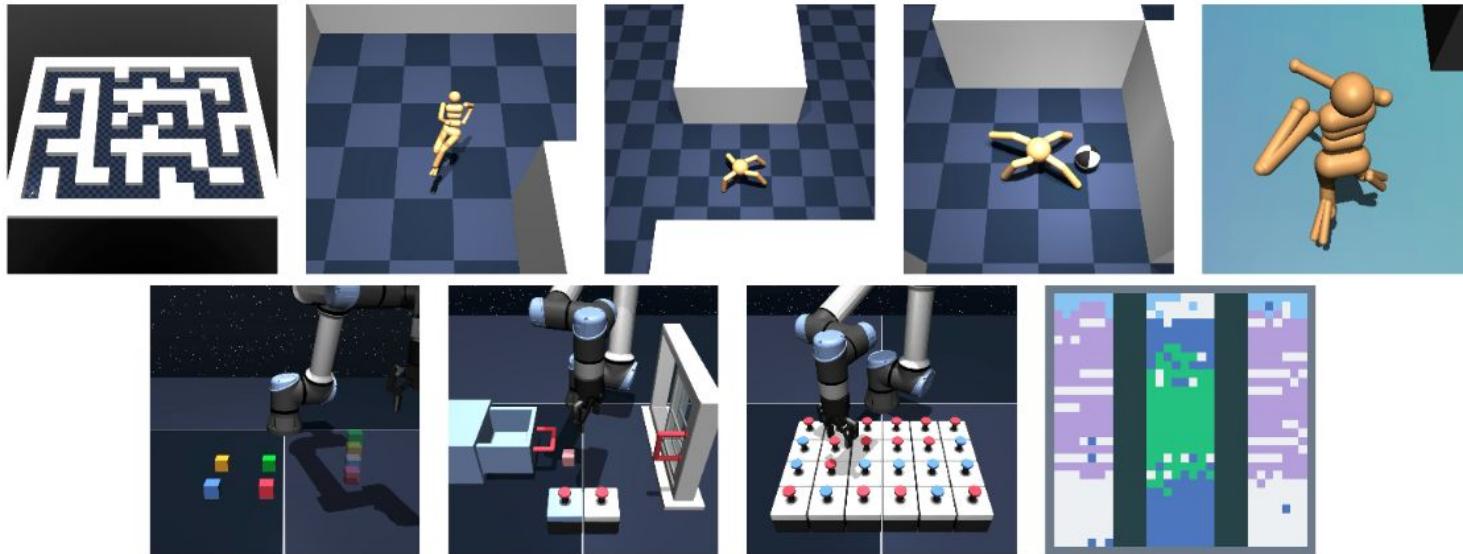
Bortkiewicz, Michał, et al. "Accelerating Goal-Conditioned RL Algorithms and Research." ICLR, 2025.

<https://github.com/MichalBortkiewicz/JaxGCRL/>

Fast Simulators: Jax GCRL



Seohong Park, Kevin Frans



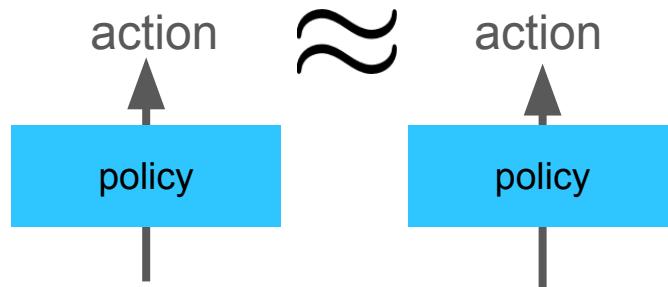
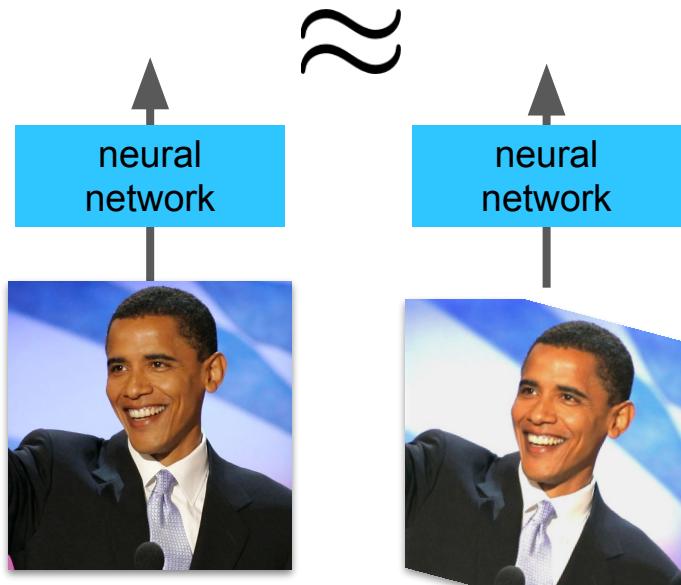
Steps towards RL agents that can do anything.

1. Fast simulators
2. **Generalization**

What is the right notion of generalization?

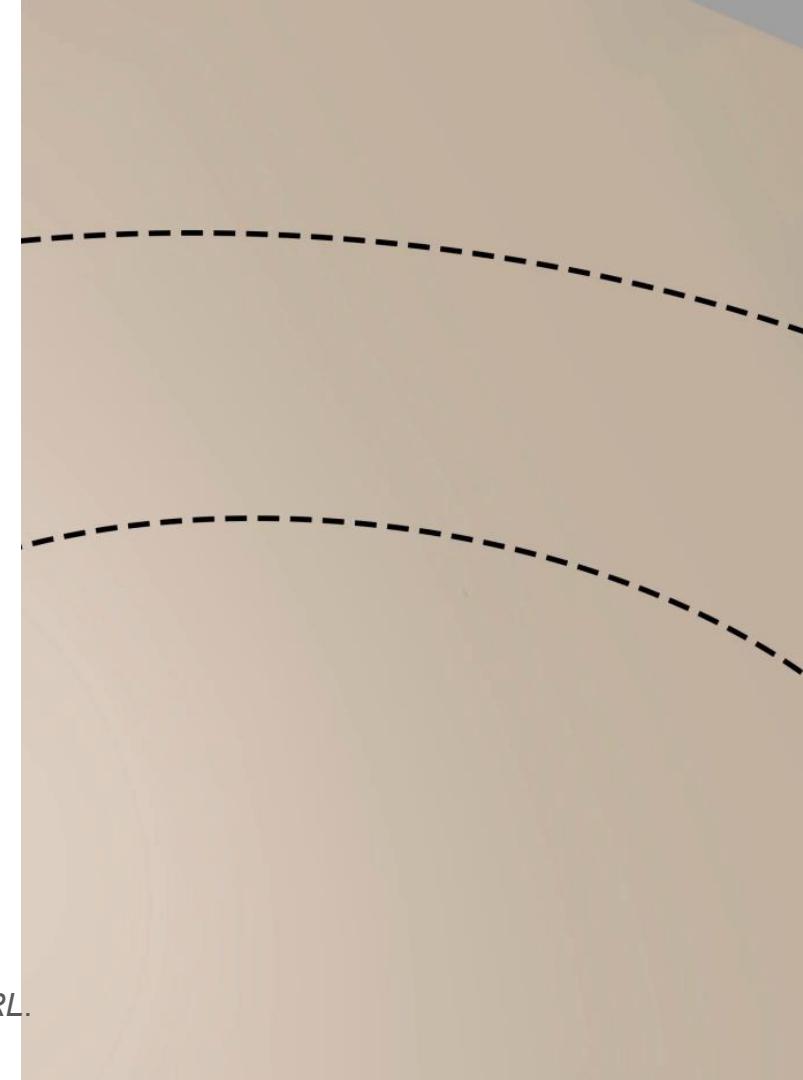


Vivek Myers, Cathy Ji


$$(s, g) \approx (s, w)$$

What is the right notion of generalization?

Horizon Generalization



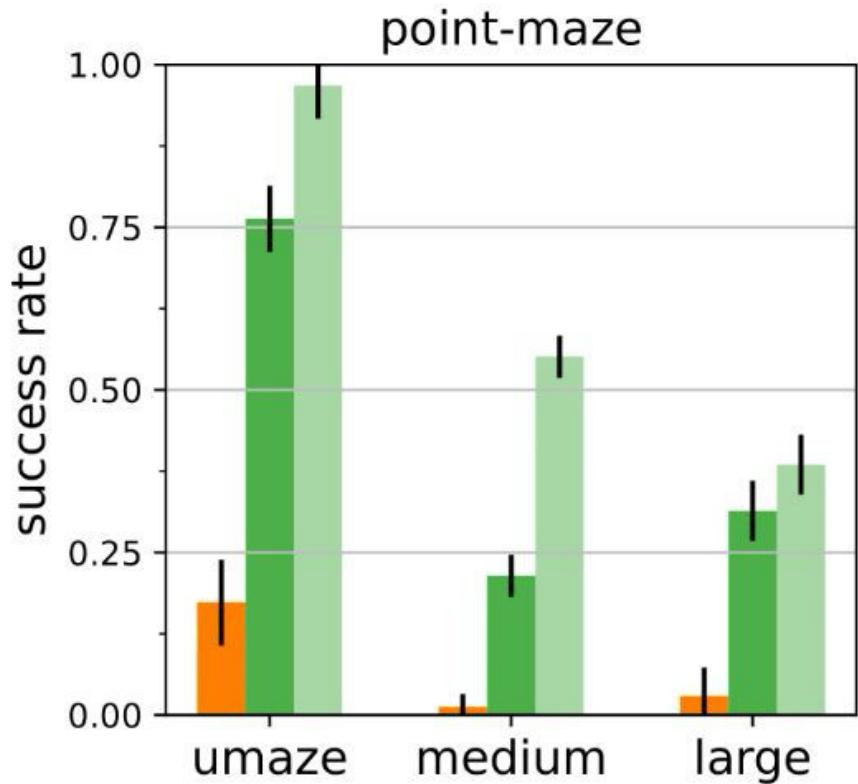
What is the right notion of generalization?

Horizon Generalization

- Data augmentation is useful



Raj Ghugare



Ghugare, Raj, et al. *Closing the Gap between TD Learning and Supervised Learning-A Generalisation Point of View*. ICLR 2024.

Steps towards RL agents that can do anything.

1. Fast simulators
2. Generalization

Self-supervised RL is generative AI
... so can we learn to do anything?

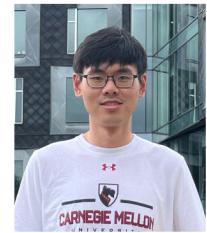
Three preliminary signs of life.

Emergent Properties in Self-Supervised RL

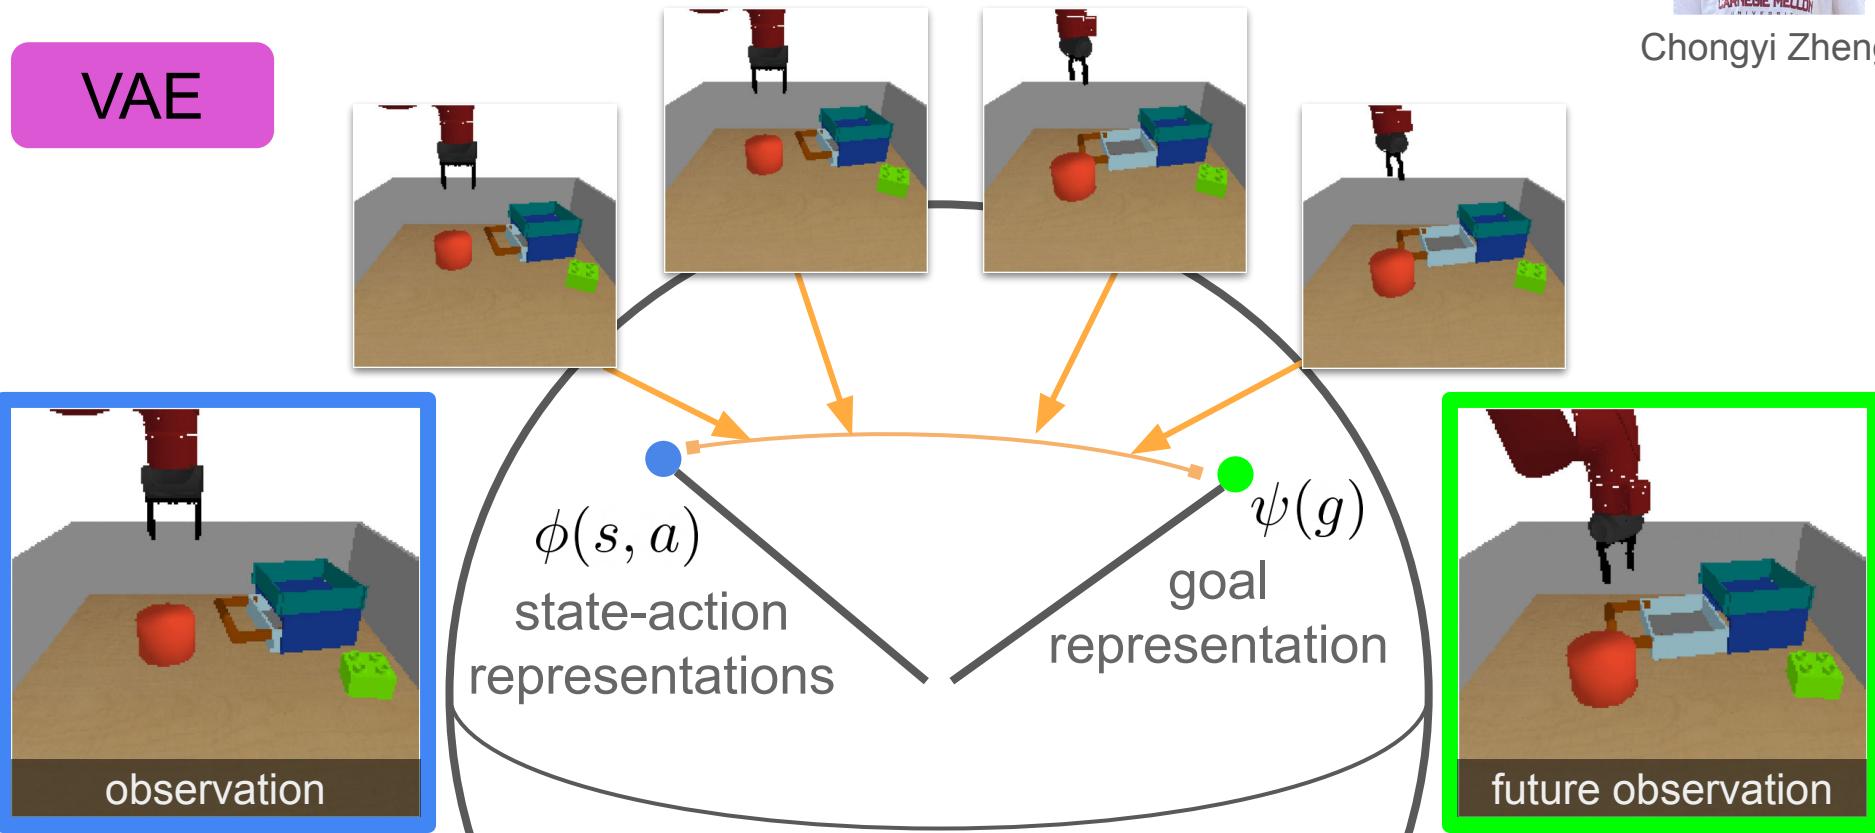
1/

VAE

Zheng, et al.
*Stabilizing Contrastive
RL: Techniques for
Robotic Goal
Reaching from Offline
Data.* ICLR, 2024



Chongyi Zheng

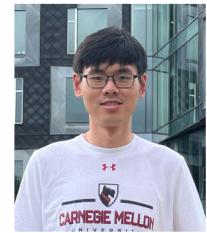


Emergent Properties in Self-Supervised RL

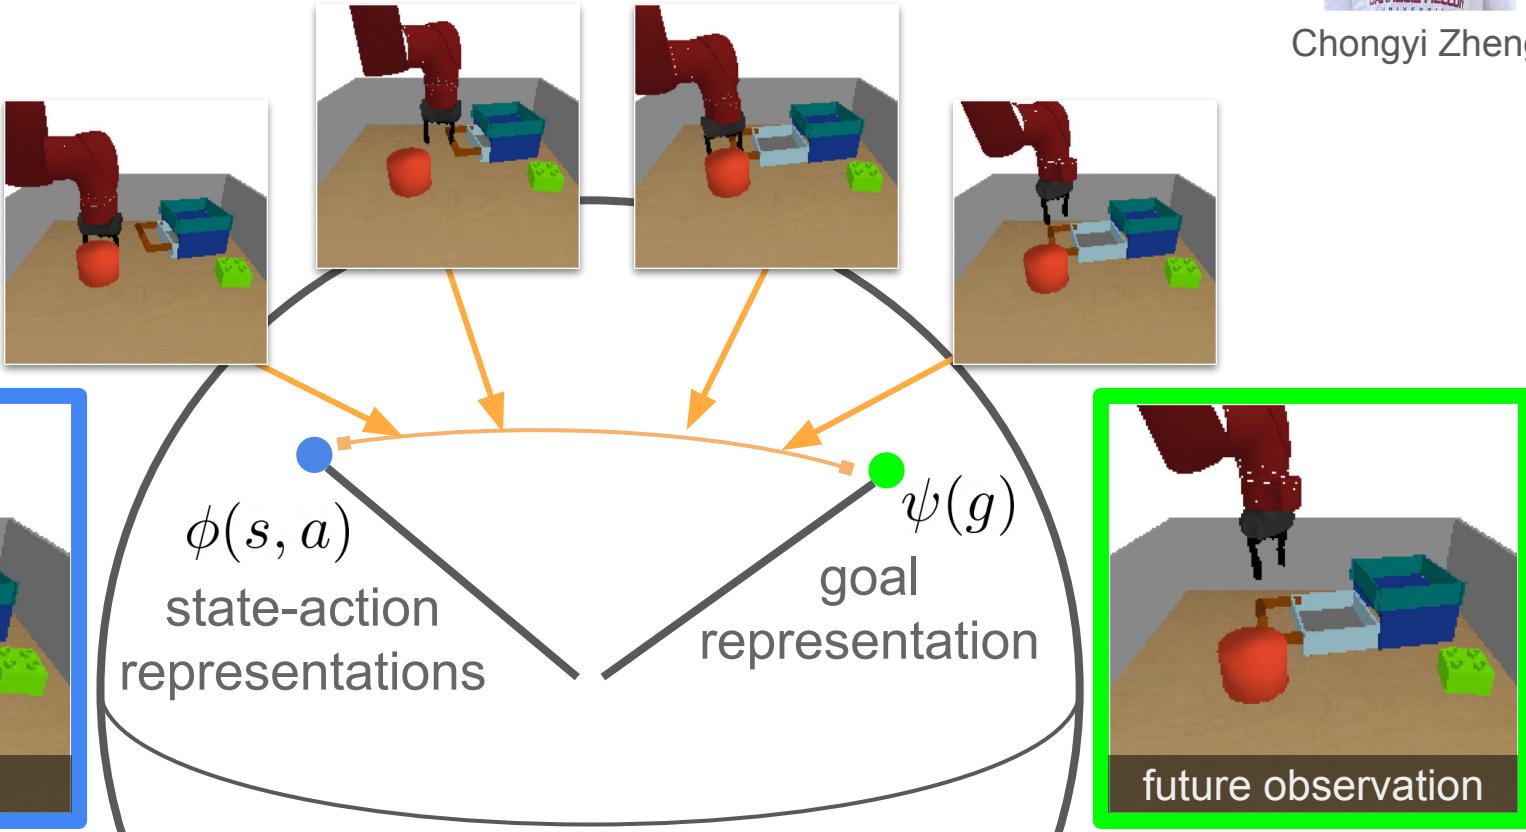
1/

Ours

Zheng, et al.
*Stabilizing Contrastive
RL: Techniques for
Robotic Goal
Reaching from Offline
Data.* ICLR, 2024



Chongyi Zheng



Emergent Properties in Self-Supervised RL

1/ Representations that Interpolate



Vivek Myers

Theorem (informal): Under some assumptions, planning over representations corresponds to inference on a Gaussian graphical model.

$$p(\psi_{1:n}) \propto \exp\left(-\frac{1}{2}\psi_{1:n}^T \Sigma^{-1} \psi_{1:n} + \eta^T \psi_{1:n}\right),$$

$$\Sigma^{-1} = \begin{pmatrix} \frac{c}{c+1} A^T A + \frac{c+1}{c} I & -A^T \\ -A & \frac{c}{c+1} A^T A + \frac{c+1}{c} I - A^T & \ddots \end{pmatrix}$$

and $\eta = \begin{pmatrix} A\psi_0 \\ 0 \\ \vdots \\ A^T \psi_{t+} \end{pmatrix}$.

Emergent Properties in Self-Supervised RL

2/



Grace Liu, Michael Tang

training goals have a *range of difficulties* (most prior methods)



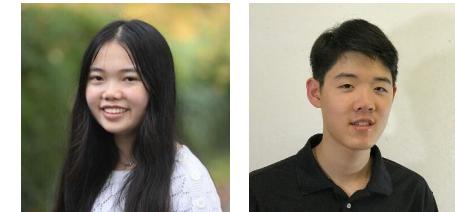
training with a *single hard goal* (ours)



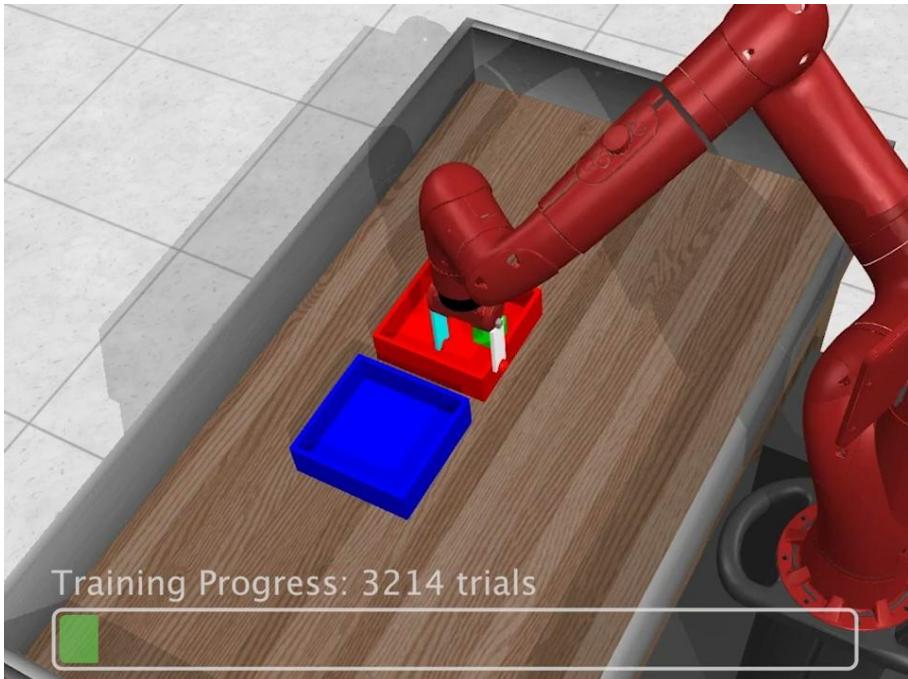
If random exploration never reaches the goal, will this learn anything?

Emergent Properties in Self-Supervised RL

2/



Grace Liu, Michael Tang



Liu, Grace, Michael Tang, and BE. *A Single Goal is All You Need: Skills and Exploration Emerge from Contrastive RL without Rewards, Demonstrations, or Subgoals*. ICLR, 2025

Emergent Properties in Self-Supervised RL

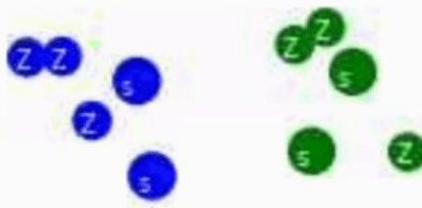
2/ Multi-Agent Exploration



Chirayu Nimonkar, Shlok Shah

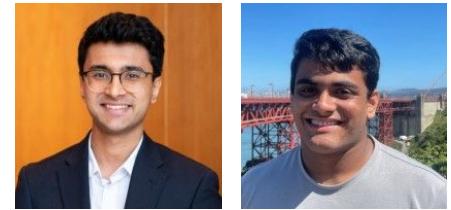


Samvelyan, Mikayel, et al. "The starcraft multi-agent challenge." arXiv preprint arXiv:1902.04043 (2019).



Emergent Properties in Self-Supervised RL

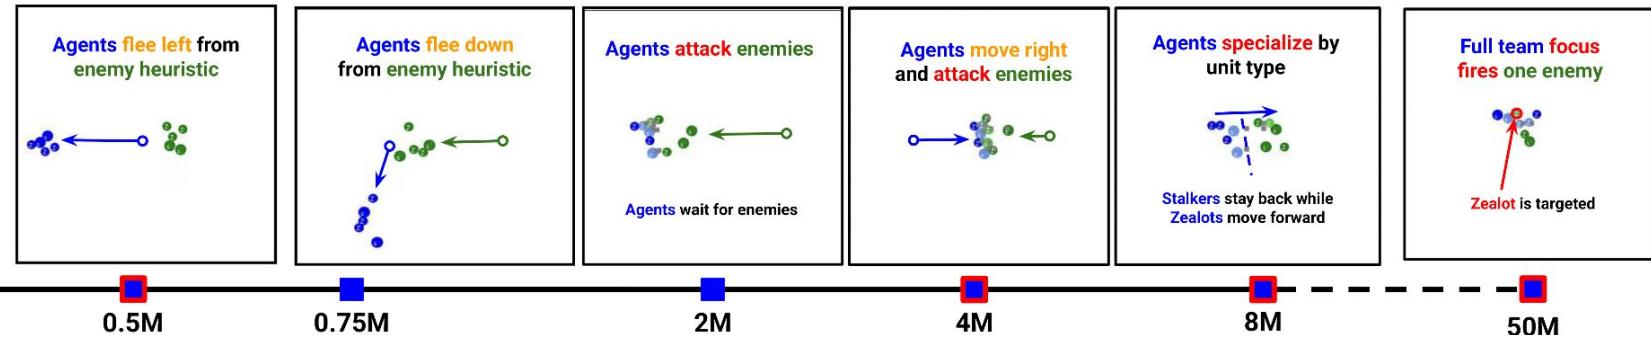
2/ Multi-Agent Exploration



Chirayu Nimonkar, Shlok Shah

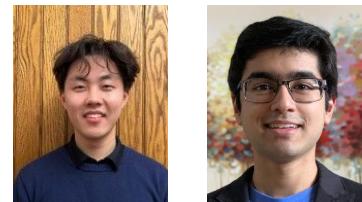
Emergent Exploration in SMAX (2s3z)

ICRL
(ours)

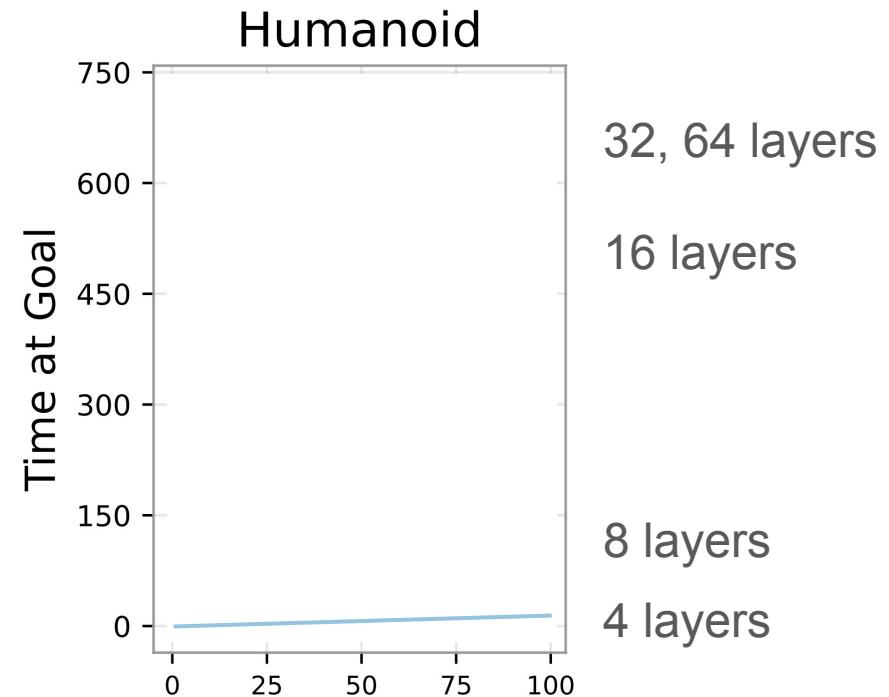
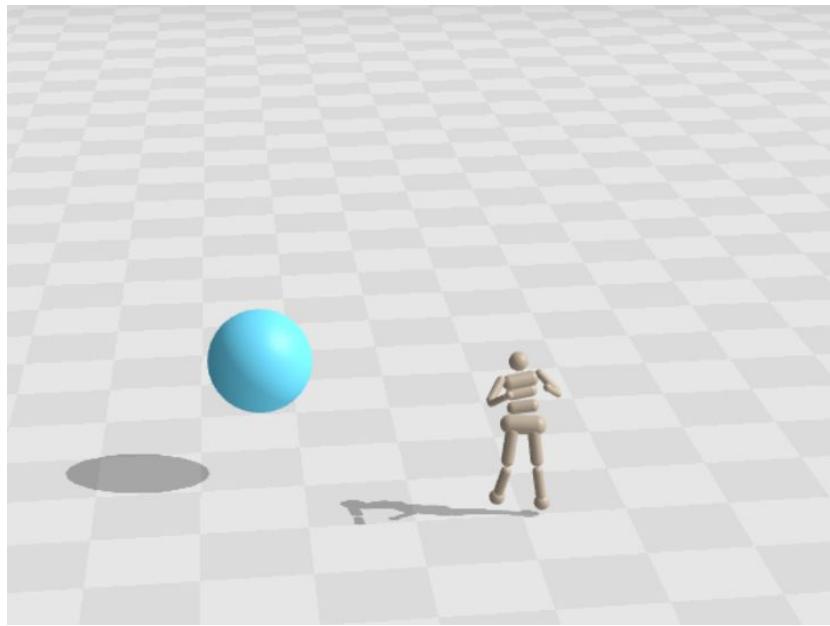


Emergent Properties in Self-Supervised RL

3/ Scale unlocks new behaviors

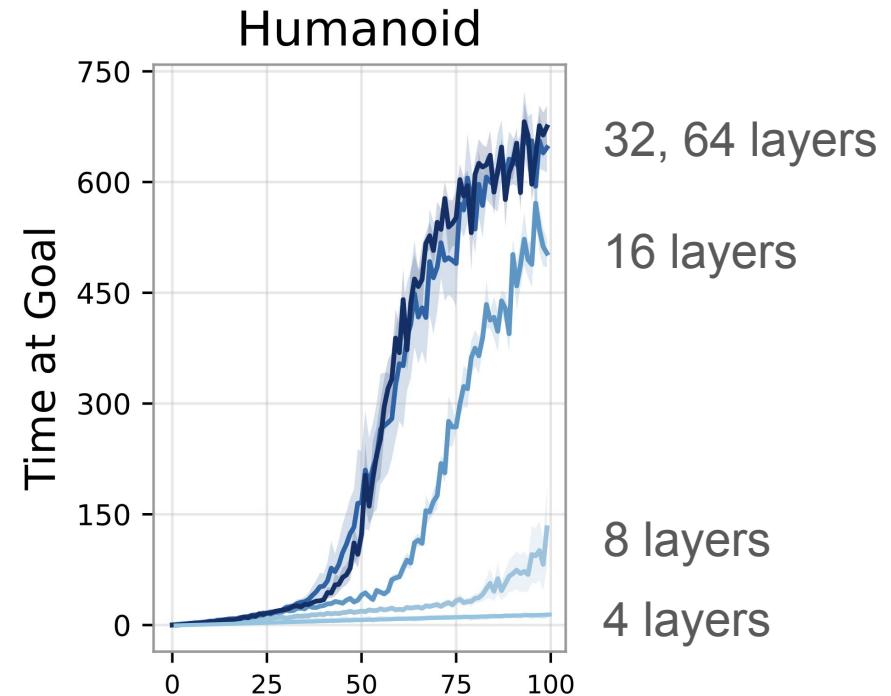
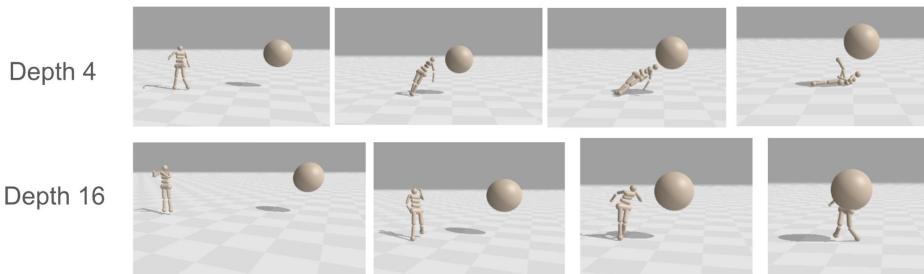
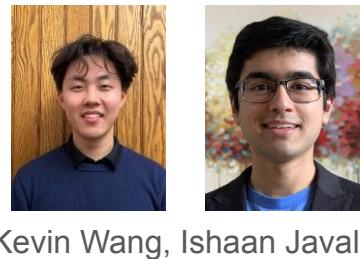


Kevin Wang, Ishaan Javali

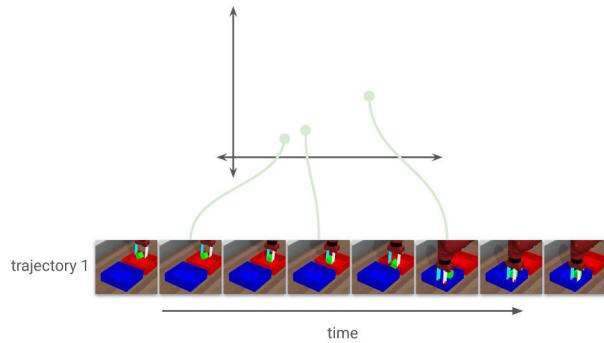


Emergent Properties in Self-Supervised RL

3/ Scale unlocks new behaviors



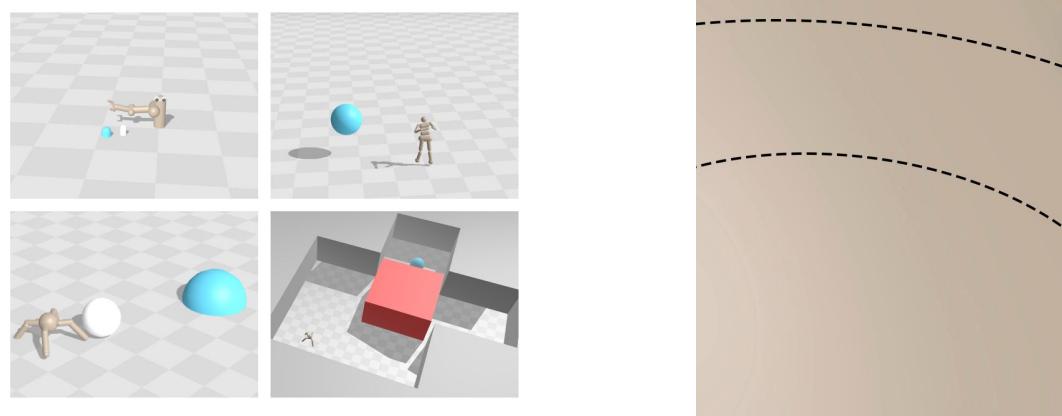
Key takeaways



1/ Self-supervised RL

Papers highlighted today:

1. BE, et al. *Contrastive learning as goal-conditioned reinforcement learning*. NeurIPS, 2022.
2. Bortkiewicz, et al. *Accelerating Goal-Conditioned RL Algorithms and Research*. ICLR, 2025.
3. Liu, et al. *A Single Goal is All You Need: Skills and Exploration Emerge from Contrastive RL without Rewards, Demonstrations, or Subgoals*. ICLR, 2025.
4. Myers, Ji, BE. *Horizon Generalization in Reinforcement Learning*. ICLR, 2025.



2/ Steps towards RL agents that can learn to do anything with minimal feedback

[Get started with research!](#)

