

# Reinforcement Learning at the Hyperscale

Jakob N. Foerster

Associate Professor (Oxford)



Foerster  
Lab for AI  
Research



St John's  
College  
University of Oxford

UNIVERSITY OF  
**OXFORD**

# Outline

- The Hyperscale Revolution of Deep RL
- Unlock #1: Faster, diverse Eval (and curriculum learning)
- Unlock #2: Theory Inspired, simpler RL
- Unlock #3: Theory Inspired, scalable Meta-RL
  - Discovering RL Algorithms
- Bonus: AI Scientist – doing e2e science.. !

Note: this is a bit of a “choose your own adventure” talk.



**FLAIR**

[Home](#)   [Members](#)   [Research](#)   [Invited Talks](#)   [Join](#)



# A tale of two revolutions: #1 The Deep Learning Revolution

## AlexNet

文 A 10 languages ▾


Article Talk

Read Edit View history Tools ▾



From Wikipedia, the free encyclopedia

**AlexNet** is the name of a convolutional neural network (CNN) architecture, designed by [Alex Krizhevsky](#) in collaboration with [Ilya Sutskever](#) and [Geoffrey Hinton](#), who was Krizhevsky's Ph.D. advisor at the University of Toronto.<sup>[1][2]</sup>

AlexNet competed in the [ImageNet Large Scale Visual Recognition Challenge](#) on September 30, 2012.<sup>[3]</sup> The network achieved a top-5 error of 15.3%, more than 10.8 percentage points lower than that of the runner up. The original paper's primary result was that the depth of the model was essential for its high performance, which was computationally expensive, but made feasible due to the utilization of [graphics processing units \(GPUs\)](#) during training.<sup>[2]</sup>



## 2: The Deep RL Revolution..



“Reinforcement Learning – An Introduction”,  
Sutton and Barto

“Playing Atari with Deep Reinforcement Learning”,  
Mnih et al, NIPS Deep Learning Workshop 2013



## Volodymyr Mnih

FOLLOW

DeepMind

Verified email at cs.toronto.edu - [Homepage](#)

Machine Learning

Looking back 12+  
years later...

TITLE

CITED BY

YEAR

[Human-level control through deep reinforcement learning](#)

30723

2015

V Mnih, K Kavukcuoglu, D Silver, AA Rusu, J Veness, MG Bellemare, ...  
Nature 518 (7540), 529-533

[Playing Atari with Deep Reinforcement Learning](#)

15140

2013

V Mnih, K Kavukcuoglu, D Silver, A Graves, I Antonoglou, D Wierstra, ...  
arXiv preprint arXiv:1312.5602



## Alex Krizhevsky

FOLLOW

[University of Toronto](#)

Verified email at cs.toronto.edu - [Homepage](#)

Machine Learning

TITLE

CITED BY

YEAR

[Imagenet classification with deep convolutional neural networks](#)

158869\*


2012

A Krizhevsky, I Sutskever, GE Hinton


Advances in neural information processing systems 25

# RL's Computational Inefficiency Problems


RL typically requires policy rollouts to happen in environments written for the CPU




# From Computational Inefficiency to Algorithmic Complexity:



(a) Batched A2C (sync step.)



(b) Batched A2C (sync traj.)



(c) IMPALA

- Difficult to keep the GPU “busy”
- Data gets stale
- Algorithmic and Engineering Complexity

“IMPALA: Importance Weighted Actor-Learner Architectures”, L Espeholt et al

# What else does this mean? Example: Hanabi!



# What else does this mean? Example: Hanabi!

*Slow*: “..Training a Rainbow agent for the *sample limited regime* of 100 million steps took approximately **seven days** using an NVIDIA V100 GPU.” [1]

- **165 steps / second on an NVIDIA V100**

*Expensive*: “.. agents for 20 billion steps. We estimate the computation took 100 CPU years for a population size of 30..” [1]

- **190 steps / CPU second**


1: “The Hanabi Challenge: A New Frontier for AI Research”, Bardt et al, <https://arxiv.org/pdf/1902.00506.pdf>



# What else does this mean? Example: Hanabi!

*Hyper-engineered*: requiring C++ for env and training loop with complicated thread handling

However, this achieves roughly 12.000 SPS on 3 GPUs.



## Take-Away:

Deep RL had lost the “Hardware Lottery”




# A Path Forward: The Second Coming of the GPU (to deep RL)




# The solution:


What if we run *everything* on the GPU instead? Wouldn't this be difficult?



# Solution: Jax!



# Solution: Jax vmap!



# Scaling Up: PureJaxRL

**NORMAL  
RL TRAINING**



**MULTIPLE  
ENVIRONMENT  
THREADS ON CPUS**

**VECTORIZED  
ENVIRONMENTS  
ON GPUs**

**END-TO-END  
VECTORIZED RL  
TRAINING ON GPUs**

## Achieving 4000x Speedups and Meta-Evolving Discoveries with PureJaxRL

### AUTHORS

Chris Lu

### AFFILIATIONS


University of Oxford,  
Foerster Lab for AI  
Research

### PUBLISHED

April 3, 2023

### FULL PAPER

arXiv



# PureJaxRL Speedups





Figure 3: CleanRL vs. Our Jax PPO on CartPole-v1 and MinAtar-Breakout. We can parallelise the agent training itself!

On CartPole-v1 we can train 2048 agents in about half the time it takes to train a single CleanRL agent!



#1:

Faster, diverse Eval  
(and curriculum discovery)

---


# KINETIX: INVESTIGATING THE TRAINING OF GENERAL AGENTS THROUGH OPEN-ENDED PHYSICS-BASED CONTROL TASKS

**Michael Matthews\***   **Michael Beukman\***   **Chris Lu**   **Jakob Foerster**  
FLAIR, University of Oxford


ICLR 2025 Oral



## Train on random environments



## Test on handmade environments



### Unified Goal

Make the **green** shape touch the **blue** shape,  
without it touching the **red** shape


Can we learn a foundation model for  
decision making?

# Interlude: Curriculum learning

## No Regrets: Investigating and Improving Regret Approximations for Curriculum Discovery

Alex Rutherford\* Michael Beukman\*  
Timon Willi Bruno Lacerda Nick Hawes Jakob Foerster

NeurIPS 2024




(a) Minigrid

(b) JaxNav


(c) Real World (Source: [18])

Figure 1: JaxNAV (b) brings UED, often tested on Minigrid (a), closer to the real world (c)



P \* ( 1- P ) (SFL) rocks

# Interlude: Curriculum learning




Kinetix Zero-Shot


$P * (1 - P)$  (SFL) rocks

# Learning to Reason at the Frontier of Learnable

Thomas Foster <sup>\*1</sup> Jakob Foerster <sup>1,2</sup>




## Fine-Tuning Results (..cause that's what you can do with a foundation model.. )



# JaxMARL: Multi-Agent RL Environments in JAX

Alexander Rutherford<sup>1\*†</sup>, Benjamin Ellis<sup>1\*†</sup>, Matteo Gallici<sup>2\*†</sup>, Jonathan Cook<sup>1\*</sup>, Andrei Lupu<sup>1\*</sup>, Garðar Ingvarsson<sup>3\*</sup>, Timon Willi<sup>1\*</sup>, Akbir Khan<sup>3</sup>, Christian Schroeder de Witt<sup>1</sup>, Alexandra Souly<sup>3</sup>, Saptarashmi Bandyopadhyay<sup>4</sup>, Mikayel Samvelyan<sup>3</sup>, Minqi Jiang<sup>3</sup>, Robert Tjarko Lange<sup>5</sup>, Shimon Whiteson<sup>1</sup>, Bruno Lacerda<sup>1</sup>, Nick Hawes<sup>1</sup>, Tim Rocktäschel<sup>3</sup>, Chris Lu<sup>1\*†</sup>, Jakob Nicolaus Foerster<sup>1</sup>

<sup>1</sup>University of Oxford <sup>2</sup>Universitat Politècnica de Catalunya <sup>3</sup>UCL <sup>4</sup>University of Maryland <sup>5</sup>Technical University Berlin




(Slides from  
Chris Lu)




# JaxMARL - So far Eight Environment Suites



MABrax




MPE



Overcooked




STORM




Antoine Bauza


ABACUSPIELE



Prisoner's  
Riddle



Coingame



SMAX


# Environment Speedups

**Table 2: Benchmark results for JAX-based MARL environments (steps-per-second) when taking random actions. All environments are significantly faster than existing CPU implementations.**


Environment	Original, 1 Env	Jax, 1 Env	Jax, 100 Envs	Jax, 10k Envs	Maximum Speedup
MPE Simple Spread	$8.34 \times 10^4$	$5.48 \times 10^3$	$5.24 \times 10^5$	$3.99 \times 10^7$	$4.78 \times 10^2$
MPE Simple Reference	$1.46 \times 10^5$	$5.24 \times 10^3$	$4.85 \times 10^5$	$3.35 \times 10^7$	$2.29 \times 10^2$
Switch Riddle	$2.69 \times 10^4$	$6.24 \times 10^3$	$7.92 \times 10^5$	$6.68 \times 10^7$	$2.48 \times 10^3$
Hanabi	$2.10 \times 10^3$	$1.36 \times 10^3$	$1.05 \times 10^5$	$5.02 \times 10^6$	$2.39 \times 10^3$
Overcooked	$1.91 \times 10^3$	$3.59 \times 10^3$	$3.04 \times 10^5$	$1.69 \times 10^7$	$8.85 \times 10^3$
MABrax Ant 4x2	$1.77 \times 10^3$	$2.70 \times 10^2$	$1.81 \times 10^4$	$7.62 \times 10^5$	$4.31 \times 10^2$
Starcraft 2s3z	$8.31 \times 10^1$	$5.37 \times 10^2$	$4.53 \times 10^4$	$2.71 \times 10^6$	$3.26 \times 10^4$
Starcraft 27m vs 30m	$2.73 \times 10^1$	$1.45 \times 10^2$	$1.12 \times 10^4$	$1.90 \times 10^5$	$6.96 \times 10^3$
STORM	—	$2.48 \times 10^3$	$1.75 \times 10^5$	$1.46 \times 10^7$	—
Coin Game	$1.97 \times 10^4$	$4.67 \times 10^3$	$4.06 \times 10^5$	$4.03 \times 10^7$	$2.05 \times 10^3$



# SMAX



...e.g.  
20k SPS on Hanabi



# See our repo and blog post for more information

## JaxMARL: Multi-Agent RL Environments and Algorithms in JAX

AUTHORS	AFFILIATIONS	PUBLISHED	FULL PAPER
Alex Rutherford, Ben Ellis, Chris Lu	University of Oxford	Nov. 17, 2023	arXiv

### Overview

We present JaxMARL, a library of multi-agent reinforcement learning (MARL) environments and algorithms based on end-to-end GPU acceleration that achieves up to 12500x speedups. The environments in JaxMARL span cooperative, competitive, and mixed games; discrete and continuous state and action spaces; and zero-shot and CTDE settings. We specifically include implementations of the Hanabi Learning Environment, Overcooked, Multi-Agent Brax, MPE, Switch Riddle, Coin Game, and Spatial-Temporal Representations of Matrix Games (STORM). **Because of JAX's hardware acceleration, our per-run training pipeline is 12500x faster than existing approaches.** We also introduce **SMAX**, a vectorised version of the popular StarCraft Multi-Agent Challenge, which removes the need to run the StarCraft II game engine. By significantly speeding up training, JaxMARL enables new research into areas such as multi-agent meta-learning, as well as significantly easing and improving evaluation in MARL. Try it out here: <https://github.com/flairox/jaxmarl>!

The screenshot shows the GitHub repository page for JaxMARL. At the top, there is a navigation bar with icons for Edit Pins, Unwatch (11), Fork (17), and Starred (192). Below the navigation bar, the repository name "JaxMARL" is displayed with a green icon. The main content area shows the "main" branch dropdown, 30 branches, and 2 tags. A list of recent commits is shown, all made by user "amacrutherford". The commits are:

Commit	Message	Time Ago
8665ca2	Create python-publish.yml	2 weeks ago
overcooked multiple seeds	overcooked multiple seeds	2 hours ago
add new doc images	add new doc images	3 days ago
Merge branch 'main' into qlearning	Merge branch 'main' into qlearning	last week
updates for flashbox	updates for flashbox	last week
name change to jaxmarl	name change to jaxmarl	3 weeks ago

On the right side of the repository page, there is a sidebar with links for About, Multi-Agent Reinforcement Learning with JAX, Readme, Apache-2.0 license, Activity, 192 stars, 11 watching, 17 forks, and Report repository. There is also a small FLAIR logo in the bottom right corner.

# JAX-LOB: A GPU-Accelerated Limit Order Book Simulator to Unlock Large Scale Reinforcement Learning for Trading

International Conference on AI in Finance: ICAIF'23  
*Best Academic Paper*

Sascha Frey, Kang Li, Peer Nagy, Silvia Sapora, Chris Lu,  
Stefan Zohren, Jakob Foerster, Anisoara Calinescu





Foerster  
Lab for AI  
Research




(Slides from  
Sasha Frey)

# The Financial Markets are a complex multi-agent System

## Limit Order Book




## Interesting Emergent Dynamics



# JAX-LOB Benefits: 75x Speed-Up Over CPU Equivalent

*Automatic parallelisation*: major benefit for **RL** since the entire training loop is GPU native. Similar speed-up expected for **Monte Carlo** methods



## Interlude 1:

# Generative AI for End-to-End Limit Order Book Modelling

A Token-Level Autoregressive Generative Model of Message Flow Using a Deep State Space Network

Peer Nagy

peer.nagy@eng.ox.ac.uk

Oxford-Man Institute of Quantitative  
Finance, University of Oxford  
UK

Kang Li

Department of Statistics, University  
of Oxford  
UK

Sascha Frey

Department of Computer Science,  
University of Oxford  
UK

Silvia Sapora

Foerster Lab for AI Research,  
University of Oxford  
UK


Stefan Zohren

Oxford-Man Institute of Quantitative  
Finance, University of Oxford &  
Man Group  
UK


Jakob Foerster

Foerster Lab for AI Research,  
University of Oxford  
UK

# A Token Level Order-Book model based on the S5 state-space model.



## Token Generation Loop



## Message Generation Loop

---

# Craftax: A Lightning-Fast Benchmark for Open-Ended Reinforcement Learning

---

Michael Matthews<sup>1</sup> Michael Beukman<sup>1</sup> Benjamin Ellis<sup>1</sup> Mikayel Samvelyan<sup>2</sup> Matthew Jackson<sup>1</sup>  
Samuel Coward<sup>1</sup> Jakob Foerster<sup>1</sup>



Michael Matthews<sup>1</sup>



Michael Beukman<sup>1</sup>



Benjamin Ellis<sup>1,2</sup>



Mikayel Samvelyan<sup>3</sup>



Matthew Jackson<sup>1,2</sup>




Samuel Coward<sup>1</sup>




Jakob Foerster<sup>1</sup>

<sup>1</sup> [FLAIR](#), University of Oxford   <sup>2</sup> [WhiRL](#), University of Oxford   <sup>3</sup> [DARK](#), University College London

## Blazing Fast



## Challenging...





#2:

GPU Acceleration + (some)

Theory =

Simplified Reinforcement

Learning

---

# Simplifying Deep Temporal Difference Learning

---

**Matteo Gallici<sup>\*</sup>**<sup>1</sup>      **Mattie Fellows<sup>\*</sup>**<sup>2</sup>

**Benjamin Ellis<sup>2</sup>**      **Bartomeu Pou<sup>1,3</sup>**      **Ivan Masmitja<sup>4</sup>**

**Jakob Nicolaus Foerster<sup>2</sup>**      **Mario Martin<sup>1</sup>**

<sup>1</sup>Universitat Politècnica de Catalunya      <sup>2</sup>University of Oxford

<sup>3</sup>Barcelona Supercomputing Center      <sup>4</sup> Institut de Ciències del Mar

{gallici,mmartin}@cs.upc.edu

{matthew.fellows,benjamin.ellis,jakob.foerster}@eng.ox.ac.uk

bartomeu.poumulet@bsc.es      masmitja@icm.csic.es

ICLR 2025 spotlight



# Part 1: Analysing and Stabilizing TD....

*Why is TD so unstable?*

It's not a gradient of anything

*How can we analyse TD?*

Using the Jacobian

*How can we stabilise TD?*

+ By introducing and LayerNorm  
+ L2 regularisation

# Theoretical insights - read the paper for details

- 1) Batchnorm layers make the policy myopic:

**Theorem 1.** *Under Assumption 1, using the BatchNorm Q-function defined in Eq. (3)*  
 $\lim_{N \rightarrow \infty} \mathbb{E}_{x_t \sim d^\mu} [\mathcal{B}[Q_\phi](x_t)] = \mathbb{E}_{r \sim P_R(x_t), x_t \sim d^\mu} [r]$  *almost surely.*

- 2) LayerNorm + L2 regularisation instead mitigates Off-Policy Instability:

**Theorem:** *Under mild regularity assumptions, using the  $k$  width LayerNorm Q-function:*



$$\frac{\left| (r(s, a) + \gamma Q_\omega(s', a') - Q_\omega(s, a)) x^\top \nabla_\omega^2 Q_\omega(s, a) x \right|}{\|x\|^2} = \mathcal{O}\left(\frac{1}{\sqrt{k}}\right)$$

almost surely for a test vector  $x$  and any  $s, a, s', a'$





Shrinks with increasing  $k$

# Tabular Validation

Baird's counterexample:  
Simple MDP with linear function  
approximator and off-policy  
sampling



# Part 2: Simplifying Temporal Difference Learning



(a) Online  $Q$ -Learning

(b) DQN

(c) Distributed DQN

(d) PQN

# No Target network, no replay – just TD Learning


---

**Algorithm 1** PQN ( $s, a, s', r, \text{target}$  are  $B$ -dim vectors)


---


- 1:  $\phi \leftarrow$  initialise regularised  $Q$ -network parameters
  - 2:  $s \leftarrow$  initial state  $s_0 \sim P_0$ ,
  - 3: **for each episode do**
  - 4:   **for each  $i \in B$  (in parallel) do**
  - 5:      $a_i \leftarrow \begin{cases} a_i \sim \text{Unif}(\mathcal{A}), & \text{with prob. } \epsilon, \\ \arg \max_{a'} Q_\phi(s_i, a'), & \text{otherwise,} \end{cases}$
  - 6:      $s'_i, r_i \leftarrow s'_i \sim P_S(s_i, a_i), r_i \sim P_R(s_i, a_i)$ ,
  - 7:      $\text{target}_i \leftarrow r_i + \gamma \mathbb{1}(\text{not terminal}) \max_{a'} Q_\phi(s'_i, a')$
  - 8:   **end for**
  - 9:    $\phi \leftarrow \phi - \alpha \nabla_\phi \| \text{StopGrad[target]} - Q_\phi(s, a) \|^2$
  - 10: **end for**
-

# The sampling regime of PQN



# Results







#3: GPU Acceleration +  
(some) Theory =

Scalable Meta - RL

# Mirror Learning Framework (ICML 2022)

---

## **Mirror Learning: A Unifying Framework of Policy Optimisation**

---

**Jakub Grudzien Kuba<sup>1</sup> Christian Schroeder de Witt<sup>1</sup> Jakob Foerster<sup>1</sup>**

# Mirror Learning: The Update Formula (Simplified)

At every iteration, maximize

$$\mathbb{E}_{a \sim \pi_{\text{new}}} [Q_{\pi_{\text{old}}}(s, a)] - \mathfrak{D}_{\pi_{\text{old}}}(\pi_{\text{new}} | s)$$

Drift Function

# Mirror Learning: Properties

**Theorem 1** (The Fundamental Theorem of Mirror Learning). *Let  $\mathfrak{D}^\nu$  be a drift,  $\mathcal{N}$  be a neighbourhood operator, and the sampling distribution  $\beta_\pi$  depend continuously on  $\pi$ . Let  $\pi_0 \in \Pi$ , and the sequence of policies  $(\pi_n)_{n=0}^\infty$  be obtained by mirror learning induced by  $\mathfrak{D}^\nu$ ,  $\mathcal{N}$ , and  $\beta_\pi$ . Then, the learned policies*

1. *Attain the strict monotonic improvement property,*

$$\eta(\pi_{n+1}) \geq \eta(\pi_n) + \mathbb{E}_{s \sim d} \left[ \frac{\nu_{\pi_n}^{\pi_{n+1}}(s)}{\beta_{\pi_n}(s)} \mathfrak{D}_{\pi_n}(\pi_{n+1}|s) \right],$$

2. *Their value functions converge to the optimal one,*


$$\lim_{n \rightarrow \infty} V_{\pi_n} = V^*,$$

3. *Their expected returns converge to the optimal return,*

$$\lim_{n \rightarrow \infty} \eta(\pi_n) = \eta^*,$$


4. *Their  $\omega$ -limit set consists of the optimal policies.*

# Space of RL Algorithms



*Figure 1.* Known RL frameworks and algorithms as points in the infinite space of *theoretically sound* mirror learning algorithms.

# Space of RL Algorithms



*Figure 1.* Known RL frameworks and algorithms as points in the infinite space of *theoretically sound* mirror learning algorithms.

---

# Discovered Policy Optimisation

---

**Chris Lu\***

FLAIR, University of Oxford  
christopher.lu@exeter.ox.ac.uk

**Jakub Grudzien Kuba\*** †

BAIR, UC Berkeley  
kuba@berkeley.edu

**Alistair Letcher**

aletcher.github.io  
ahp.letcher@gmail.com

**Luke Metz**

Google Brain  
Luke.s.metz@gmail.com

**Christian Schroeder de Witt**

FLAIR, University of Oxford  
cs@robots.ox.ac.uk

**Jakob Foerster**

FLAIR, University of Oxford  
jakob.foerster@eng.ox.ac.uk

NeurIPS 2022



# A Meta Learning Approach


At every iteration, maximize

$$\mathbb{E}_{a \sim \pi_{\text{new}}} [Q_{\pi_{\text{old}}}(s, a)] - \mathfrak{D}_{\pi_{\text{old}}}(\pi_{\text{new}} | s)$$

Drift Function


Parameterise as a neural network,  $\phi$ , compliant with theory

## How to estimate Meta-Gradients?



$$\nabla_{\phi} \mathbb{E}_{\epsilon \sim N(0, I)} [F(\phi + \sigma \epsilon)] = \frac{1}{\sigma} \mathbb{E}_{\epsilon \sim N(0, I)} [F(\phi + \sigma \epsilon) \epsilon],$$


# Learned Policy Optimisation (LPO) Performance



Note: All runs use BRAX default Hyperparameters for PPO

# Visualizing Objective Functions: PPO


$$A_{\pi}(s, a)$$



$$\mathbf{r}(\bar{\pi}) \triangleq \bar{\pi}(a|s)/\pi(a|s)$$



# Visualizing Objective Functions: PPO

$$A_{\pi}(s, a)$$




$$\mathbf{r}(\bar{\pi}) \triangleq \bar{\pi}(a|s)/\pi(a|s)$$

# Visualizing Objective Functions: LPO




# Interpreting LPO




1. “Rollback” for Negative Advantage
2. Cautious Optimism for Positive Advantage
3. Implicit Entropy Maximisation
4. Secondary Features

# Interpreting LPO




1. Rollback for Negative Advantage
2. Cautious Optimism for Positive Advantage
3. Implicit Entropy Maximisation
4. Secondary Features

# Interpreting LPO



1. Rollback for Negative Advantage
2. Cautious Optimism for Positive Advantage
3. Implicit Entropy Maximisation
4. Secondary Features


# Interpreting LPO




1. Rollback for Negative Advantage
2. Cautious Optimism for Positive Advantage
3. Implicit Entropy Maximisation
4. Secondary Features

# Discovered Policy Optimisation (DPO)

$$f(r, A) = \begin{cases} \text{ReLU}\left((r - 1)A - \alpha \tanh\left((r - 1)A/\alpha\right)\right) & A \geq 0 \\ \text{ReLU}\left(\log(r)A - \beta \tanh\left(\log(r)A/\beta\right)\right) & A < 0 \end{cases}$$




# Discovered Policy Optimisation Performance




Note: All runs use BRAX default Hyperparameters for PPO


# Discovered Policy Optimisation: Far OOD Performance




(a) Breakout-Minatar



(b) Asterix-Minatar



(c) Freeway-Minatar



(d) SpaceInvaders-Minatar


# Discovery: DPO

**1. Picked**  $\phi$  that corresponds to a core part of hand-crafted RL objective functions. Parameterised based on theory

**2. Meta-Optimized**  $\phi$  using evolution and PureJaxRL

**3. Analyzed**  $\phi$  to gain insight into policy optimization and create Discovered Policy Optimisation

# Temporally-Adaptive LPO (ICLR 2024)



# We get much better performance!




Figure 2: **TA-LPO leverages lifetime information and generalizes to a wide range of environments.** Results of TA-LPO, LPO and PPO on the Brax and MinAtar suites across three seeds. TA-LPO was only meta-trained on SpaceInvaders-MinAtar. We provide complete training curves in Appendix B.

Further example work from FLAIR at the *Hyperscale*..

## 1. Reward Functions

- a. Sapora, Silvia, et al. "EvIL: Evolution Strategies for Generalisable Imitation Learning." *ICML 2024*

## 2. Expert Datasets

- a. Lupu, Andrei, et al. "Behaviour Distillation." *ICLR 2024*

## 3. High-Dimensional Opponent Shaping

- a. Khan, Akbir, et al. "Scaling Opponent Shaping to High Dimensional Games." *AAMAS 2023*

## 4. Better Environments

- a. Matthews, Michael et al. "Craftax" *ICML 2024*
- b. Frey, Sasha et al, "JAX-LOB" *ICAIF 2023*

## 5. Synthetic Environments

- a. Lu, Chris et al. "Adversarial Cheap Talk." *ICML 2023*
- b. Liesen, Jarek et al. "Discovering Minimal Reinforcement Learning Environments"

## 6. Learned RL Optimizers

- a. Goldie, Alexander et al. "Can Learned Optimization Make Reinforcement Learning Less Difficult?" *ICML 2024 AutoRL Workshop (Spotlight)*

All of these experiments would have taken *years* to run in the old paradigm...



# Conclusion

- We are in the middle of a revolution for deep RL: *end-to-end GPU acceleration*
- Speed-ups of many orders of magnitude are possible
- We have implemented a number of environments in JAX
- This unlocks theory-inspired, simpler, highly performant algorithms
- It also pushes the frontier for meta-RL
- All code and many tools are open source on our github
- Thanks for listening!

Open questions:

- RL at the hyperscale in the Age of (Agentic) LLMs?
- Breaking out of the (JAX) Box?

Thanks for listening!

More info and code [@j\\_foerst](https://github.com/jfoerster) and [www.foersterlab.com](http://www.foersterlab.com)

